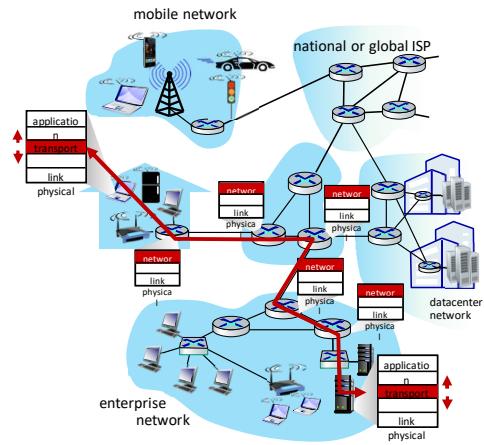
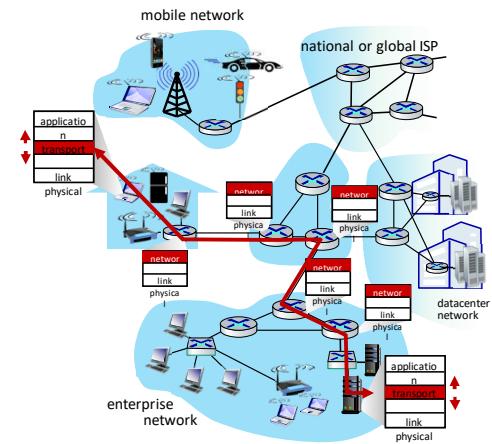


Network layer



Network-layer services and protocols

- transport segment from sending to receiving host
 - **sender:** encapsulates segments into datagrams, passes to link layer
 - **receiver:** delivers segments to transport layer protocol
- network layer protocols in ***every Internet device***: hosts, routers
- **routers:**
 - examines header fields in all IP datagrams passing through it
 - moves datagrams from input ports to output ports to transfer datagrams along end-end path



network-layer functions:

- **forwarding**: move packets from a router's input link to appropriate router output link
- **routing**: determine route taken by packets from source to destination
 - *routing algorithms*

analogy: taking a trip

- **forwarding**: process of getting through single interchange
- **routing**: process of planning trip from source to destination



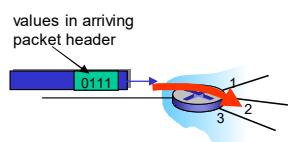
forwarding



routing

Data plane:

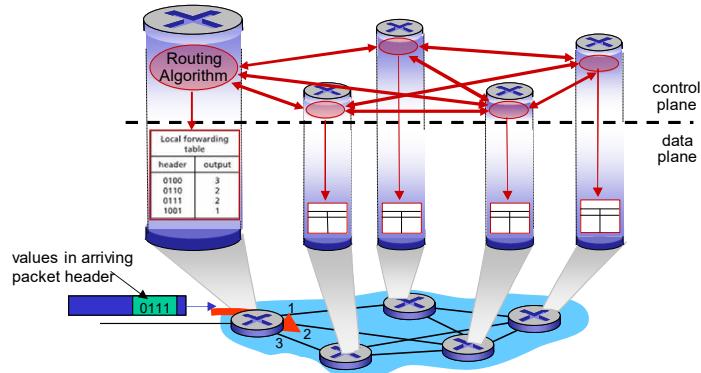
- *local*, per-router function
- determines how datagram arriving on router input port is forwarded to router output port

**Control plane**

- *network-wide* logic
- determines how datagram is routed among routers along end-end path from source host to destination host
- two control-plane approaches:
 - *traditional routing algorithms*: implemented in routers
 - *software-defined networking (SDN)*: implemented in (remote) servers

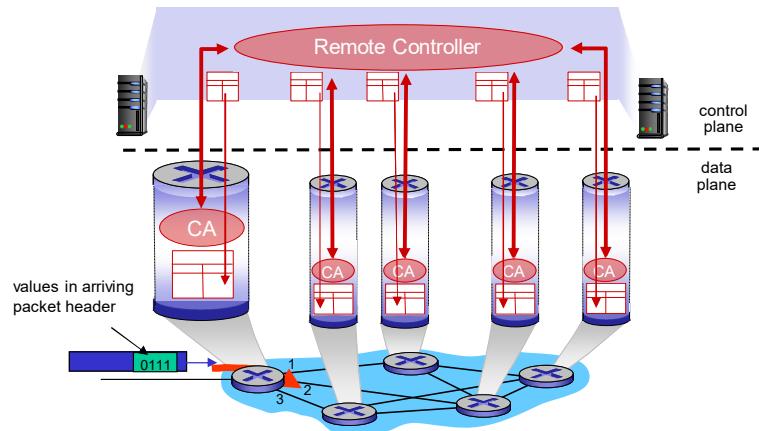
Per-router control plane

Individual routing algorithm components *in each and every router* interact in the control plane



Software-Defined Networking (SDN) control plane

Remote controller computes, installs forwarding tables in routers



Q: What *service model* for “channel” transporting datagrams from sender to receiver?

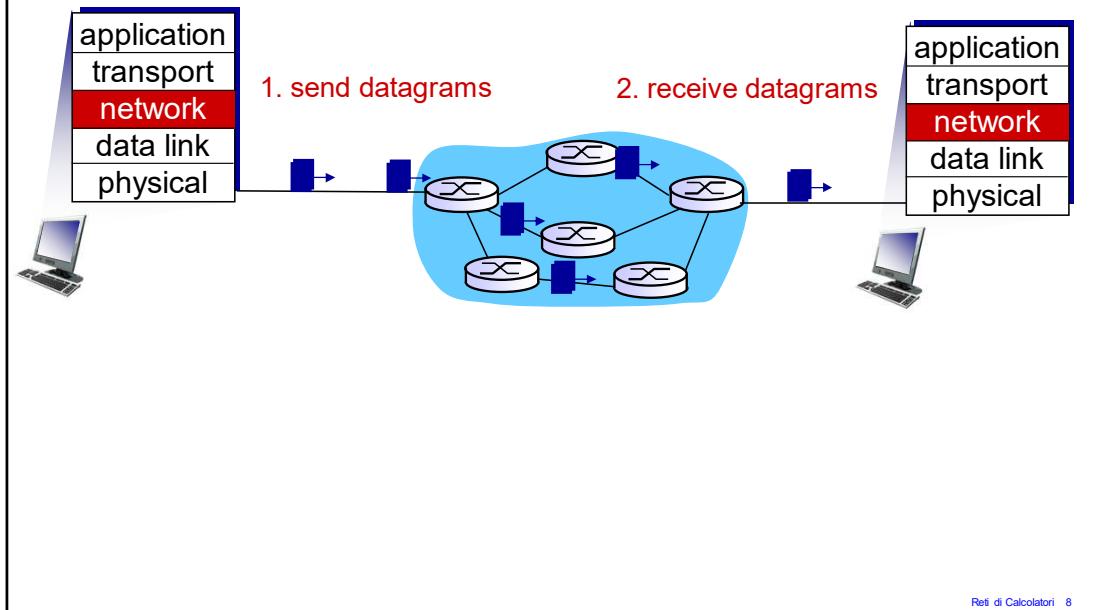
example services for *individual* datagrams:

- guaranteed delivery
- guaranteed delivery with less than 40 msec delay

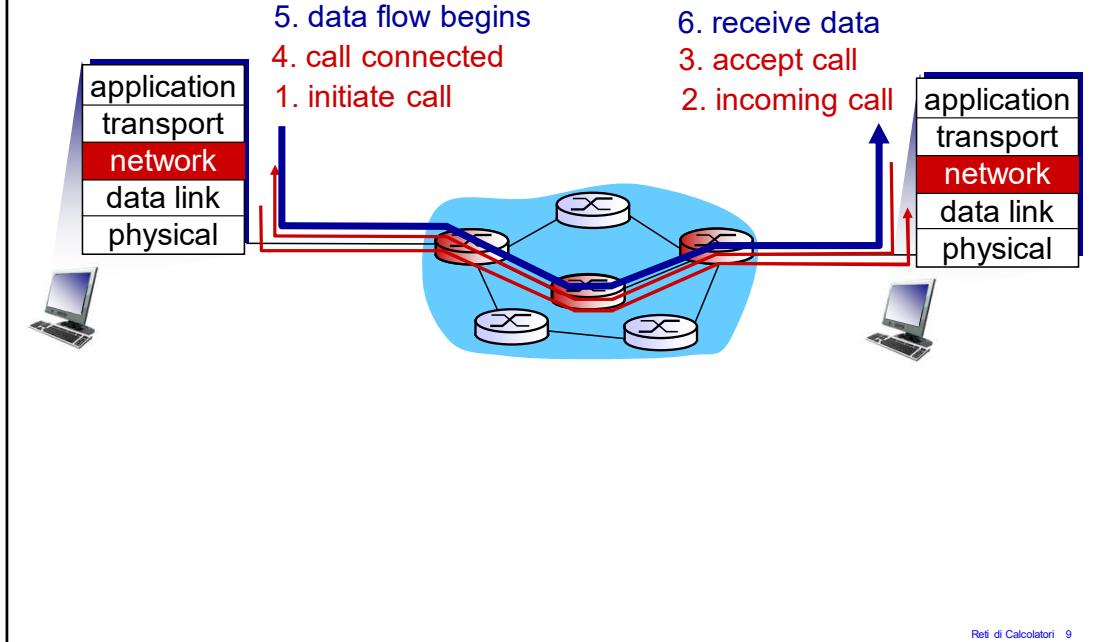
example services for a *flow* of datagrams:

- in-order datagram delivery
- guaranteed minimum bandwidth to flow
- restrictions on changes in inter-packet spacing

Datagram service



Reti di Calcolatori 8



	Datagram	Virtual circuit
Circuit creation	No	Yes
State information	Address in every packet	VC address
Routing	No information	Every VC must be identified
Effects of router faults	No	Fault for every VC involved
Congestion control	Complex	Easy

Network Architecture	Service Model	Quality of Service (QoS) Guarantees ?			
		Bandwidth	Loss	Order	Timing
Internet	best effort	none	no	no	no

Internet “best effort” service model

No guarantees on:

- 1) successful datagram delivery to destination
- 2) timing or order of delivery
- 3) bandwidth available to end-end flow

Network Architecture	Service Model	Quality of Service (QoS) Guarantees ?			
		Bandwidth	Loss	Order	Timing
Internet	best effort	none	no	no	no
ATM	Constant Bit Rate	Constant rate	yes	yes	yes
ATM	Available Bit Rate	Guaranteed min	no	yes	no
Internet	Intserv Guaranteed (RFC 1633)	yes	yes	yes	yes
Internet	Diffserv (RFC 2475)	possible	possibly	possibly	no

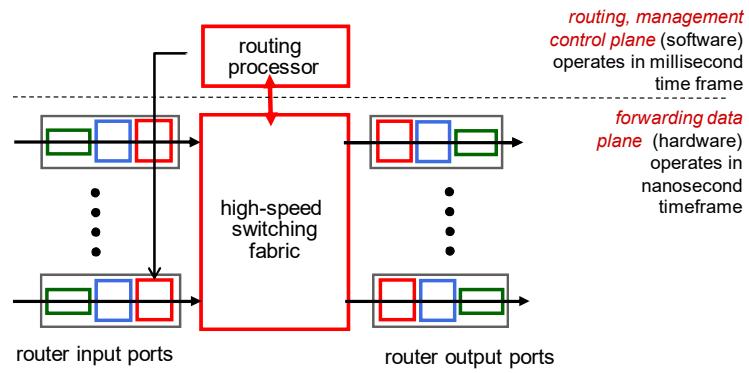
- simplicity of mechanism has allowed Internet to be widely deployed adopted
- sufficient provisioning of bandwidth allows performance of real-time applications (e.g., interactive voice, video) to be “good enough” for “most of the time”
- replicated, application-layer distributed services (datacenters, content distribution networks) connecting close to clients’ networks, allow services to be provided from multiple locations
- congestion control of “elastic” services helps

It's hard to argue with success of best-effort service model

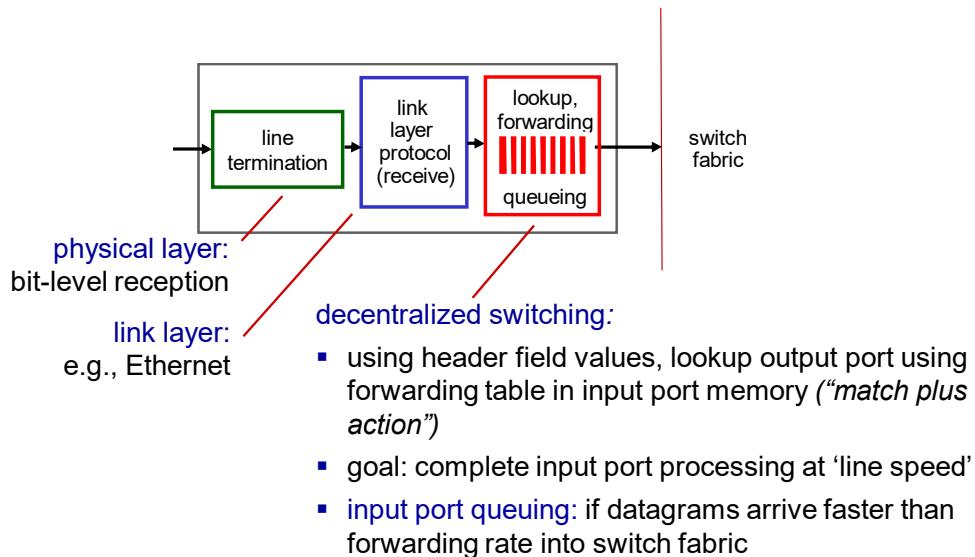
- Network layer: overview
 - data plane
 - control plane
- What's inside a router
 - input ports, switching, output ports
 - buffer management, scheduling
- IP: the Internet Protocol
 - datagram format
 - addressing
 - network address translation
 - IPv6



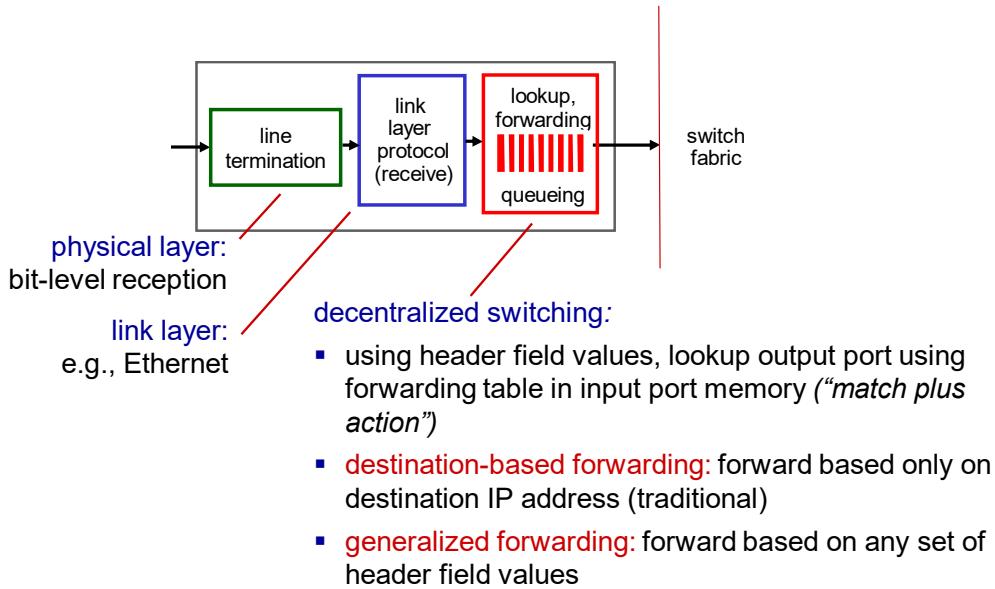
High-level view of generic router architecture:



Input port functions



Input port functions



forwarding table

Destination Address Range	Link Interface
11001000 00010111 00010000 00000000 through 11001000 00010111 00010111 11111111	0
11001000 00010111 00011000 00000000 through 11001000 00010111 00011000 11111111	1
11001000 00010111 00011001 00000000 through 11001000 00010111 00011111 11111111	2
otherwise	3

Q: but what happens if ranges don't divide up so nicely?

Destination-based forwarding

forwarding table

Destination Address Range	Link Interface
11001000 00010111 00010000 00000000 through 11001000 00010111 00010111 11111111	0
11001000 00010111 00010000 00000100 through 11001000 00010111 00010000 00000111	3
11001000 00010111 00011000 00000000 through 11001000 00010111 00011000 11111111	1
11001000 00010111 00011001 00000000 through 11001000 00010111 00011111 11111111	2
otherwise	3

longest prefix match

when looking for forwarding table entry for given destination address, use *longest* address prefix that matches destination address.

Destination Address Range	Link interface
11001000 00010111 00010*** *****	0
11001000 00010111 00011000 *****	1
11001000 00010111 00011*** *****	2
otherwise	3

examples:

11001000 00010111 00010110 10100001 which interface?

11001000 00010111 00011000 10101010 which interface?

longest prefix match

when looking for forwarding table entry for given destination address, use **longest** address prefix that matches destination address.

Destination Address Range	Link interface
11001000 00010111 00010*** ****	0
11001000 00010111 00011000 ****	1
11001000 match! 00011*** ****	2
otherwise	3

examples:

11001000 00010111 00010110 10100001 which interface?

11001000 00010111 00011000 10101010 which interface?

longest prefix match

when looking for forwarding table entry for given destination address, use *longest* address prefix that matches destination address.

Destination Address Range	Link interface
11001000 00010111 00010*** *****	0
11001000 00010111 00011000 *****	1
11001000 00010111 00011*** *****	2
otherwise	3

examples:

11001000 00010111 00010110 10100001 which interface?

match!

11001000 00010111 00011000 10101010 which interface?

longest prefix match

when looking for forwarding table entry for given destination address, use *longest* address prefix that matches destination address.

Destination Address Range	Link interface
11001000 00010111 00010*** *****	0
11001000 00010111 00011000 *****	1
11001000 00010111 00011*** *****	2
otherwise match!	3

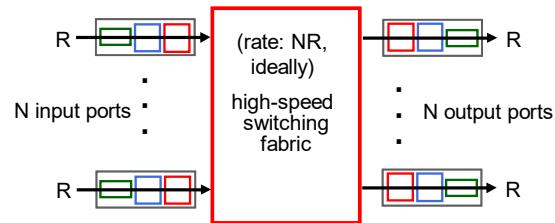
examples:

11001000 00010111 00010110 10100001 which interface?

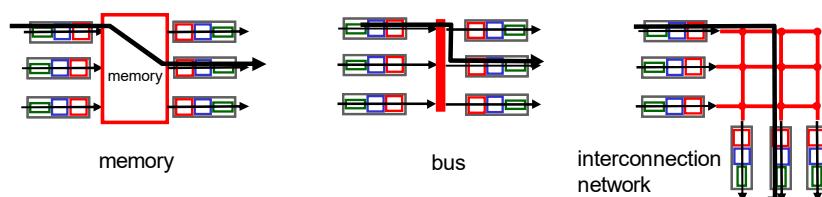
11001000 00010111 00011000 10101010 which interface?

- we'll see *why* longest prefix matching is used shortly, when we study addressing
- longest prefix matching: often performed using ternary content addressable memories (TCAMs)
 - **content addressable**: present address to TCAM: retrieve address in one clock cycle, regardless of table size
 - Cisco Catalyst: ~1M routing table entries in TCAM

- transfer packet from input link to appropriate output link
- **switching rate:** rate at which packets can be transferred from inputs to outputs
 - often measured as multiple of input/output line rate
 - N inputs: switching rate N times line rate desirable

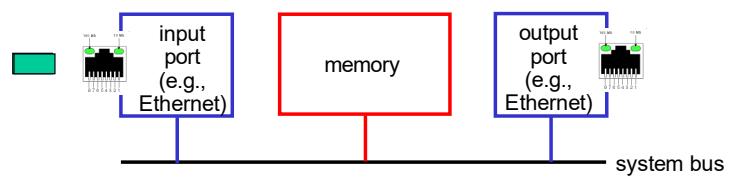


- transfer packet from input link to appropriate output link
- **switching rate:** rate at which packets can be transferred from inputs to outputs
 - often measured as multiple of input/output line rate
 - N inputs: switching rate N times line rate desirable
- three major types of switching fabrics:

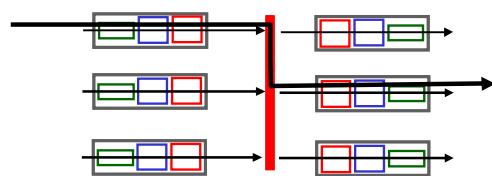


first generation routers:

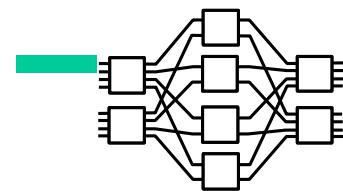
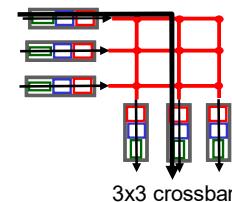
- traditional computers with switching under direct control of CPU
- packet copied to system's memory
- speed limited by memory bandwidth (2 bus crossings per datagram)



- datagram from input port memory to output port memory via a shared bus
- **bus contention:** switching speed limited by bus bandwidth
- 32 Gbps bus, Cisco 5600: sufficient speed for access routers

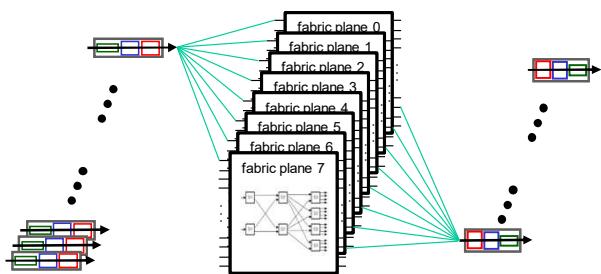


- Crossbar, Clos networks, other interconnection nets initially developed to connect processors in multiprocessor
- **multistage switch:** nxn switch from multiple stages of smaller switches
- **exploiting parallelism:**
 - fragment datagram into fixed length cells on entry
 - switch cells through the fabric, reassemble datagram at exit

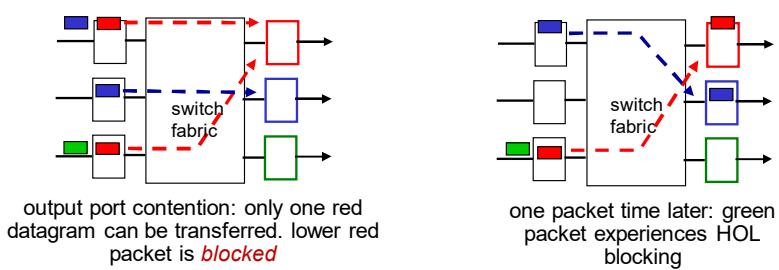


- scaling, using multiple switching “planes” in parallel:
 - speedup, scaleup via parallelism

- Cisco CRS router:
 - basic unit: 8 switching planes
 - each plane: 3-stage interconnection network
 - up to 100's Tbps switching capacity



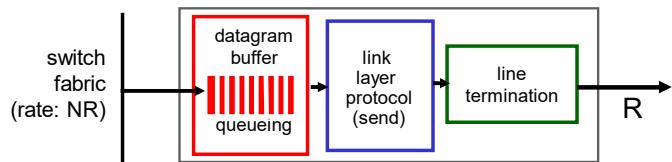
- If switch fabric slower than input ports combined -> queueing may occur at input queues
 - queueing delay and loss due to input buffer overflow!
- **Head-of-the-Line (HOL) blocking:** queued datagram at front of queue prevents others in queue from moving forward



Output port queuing

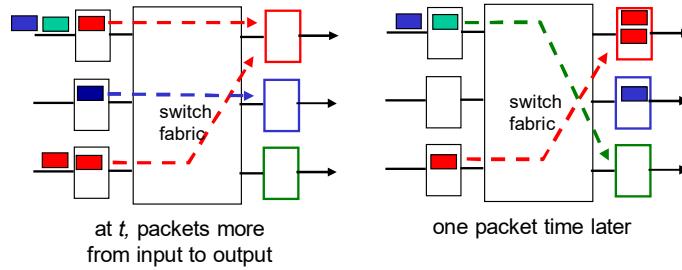


This is a really important slide



- **Buffering** required when datagrams arrive from fabric faster than link transmission rate. **Drop policy:** which datagrams to drop if no free buffers?
 - Datagrams can be lost due to congestion, lack of buffers
- **Scheduling discipline** chooses among queued datagrams for transmission
 - Priority scheduling – who gets best performance, network neutrality

Output port queuing



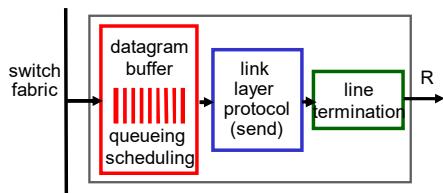
- buffering when arrival rate via switch exceeds output line speed
- **queueing (delay) and loss due to output port buffer overflow!**

- RFC 3439 rule of thumb: average buffering equal to “typical” RTT (say 250 msec) times link capacity C
 - e.g., C = 10 Gbps link: 2.5 Gbit buffer

- more recent recommendation: with N flows, buffering equal to

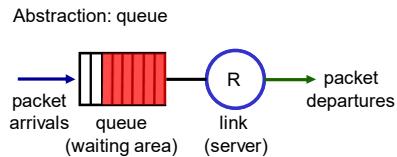
$$\frac{\text{RTT} \cdot C}{\sqrt{N}}$$

- but *too* much buffering can increase delays (particularly in home routers)
 - long RTTs: poor performance for realtime apps, sluggish TCP response
 - recall delay-based congestion control: “keep bottleneck link just full enough (busy) but no fuller”



buffer management:

- **drop:** which packet to add, drop when buffers are full
 - **tail drop:** drop arriving packet
 - **priority:** drop/remove on priority basis
- **marking:** which packets to mark to signal congestion (ECN, RED)



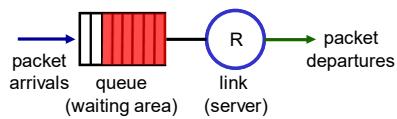
packet scheduling: deciding which packet to send next on link

- first come, first served
- priority
- round robin
- weighted fair queueing

FCFS: packets transmitted in order of arrival to output port

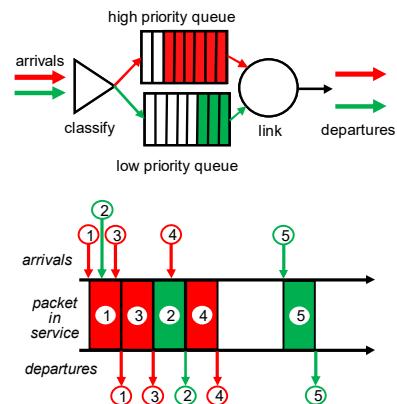
- also known as: First-in-first-out (FIFO)
- real world examples?

Abstraction: queue



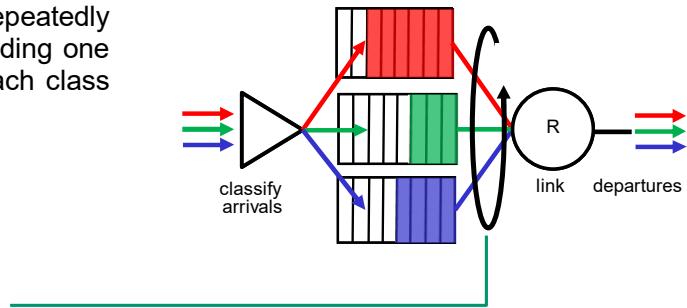
Priority scheduling:

- arriving traffic classified, queued by class
 - any header fields can be used for classification
- send packet from highest priority queue that has buffered packets
 - FCFS within priority class



Round Robin (RR) scheduling:

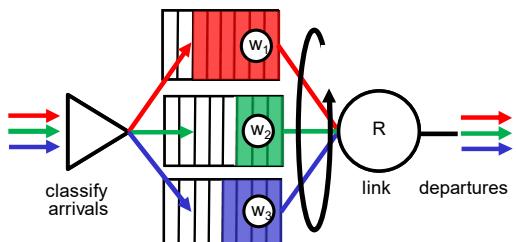
- arriving traffic classified, queued by class
 - any header fields can be used for classification
- server cyclically, repeatedly scans class queues, sending one complete packet from each class (if available) in turn

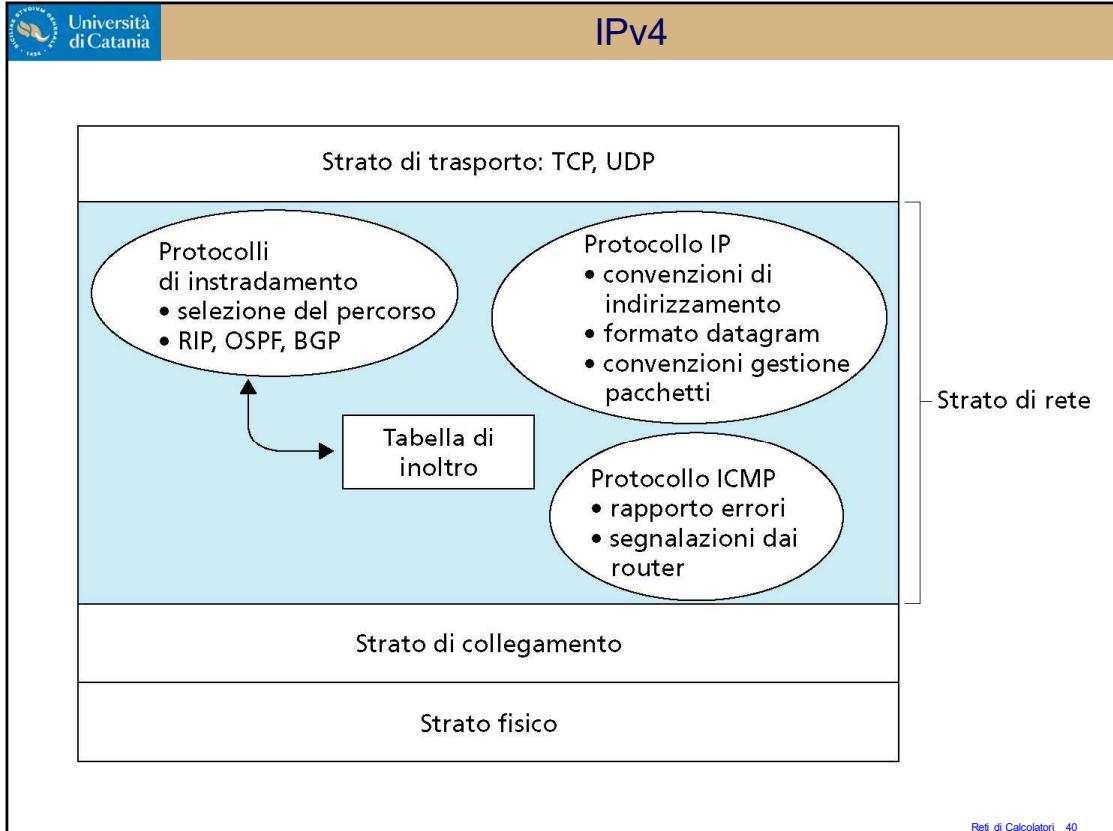


Weighted Fair Queueing (WFQ):

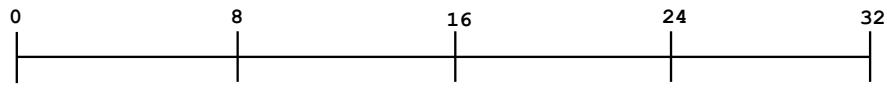
- generalized Round Robin
- each class, i , has weight, w_i , and gets weighted amount of service in each cycle:

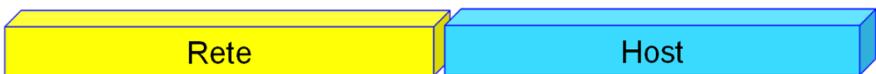
$$\frac{w_i}{S_j w_j}$$
- minimum bandwidth guarantee (per-traffic-class)





Indirizzi IPv4


A: 0...

B: 10...

C: 110...


Classe	Da	A
A	1.0.0.0	127.255.255.255
B	128.0.0.0	191.255.255.255
C	192.0.0.0	223.255.255.255
D (multicast)	224.0.0.0	239.255.255.255
E (riservati)	240.0.0.0	

0.0.0.0	This host
00000....xxxxxxxxxx	Host nella subnet
255.255.255.255	Broadcast
127.0.0.0 – 127.255.255.255	Loopback
10.0.0.0 – 10.255.255.255	IP Privati
169.254.0.0 – 169.254.255.255	Zero conf. net. APIPA
172.16.0.0 – 172.31.255.255	IP Privati
192.168.0.0 – 192.168.255.255	IP Privati

Inizialmente era effettuato da **IANA** (Internet Assigned Numbers Authority) (www.iana.org).

Nel 1992 sono state create le **RIR** (Regional Internet Registry), per gestire gli indirizzi e i nomi di dominio per ciascuna area assegnata.

Nel 1998 è stata fondata un'organizzazione no-profit denominata **ICANN** (Internet Corporation for Assigned Network and Numbers, www.icann.org) sotto il controllo del ministero del commercio USA per coordinare tutte le RIR del mondo.

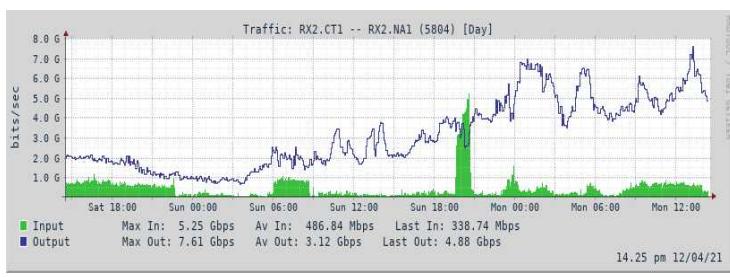
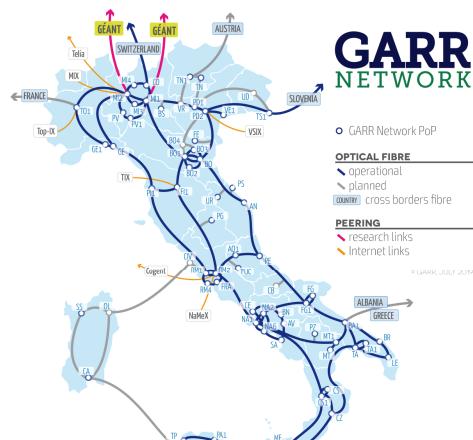
Oggi esistono cinque RIR:

- **APNIC** (Asia Pacific Network Information Center): www.apnic.net.
- **ARIN** (American Registry for Internet Numbers): www.arin.net.
- **LACNIC** (Latin American and Caribbean Internet Addresses Registry): www.lacnic.net.
- **RIPE NCC** (Réseaux IP Européens Network Coordination Centre): www.ripe.net.
- **AFRINIC** (African Regional Internet Registry): www.afrinic.org.

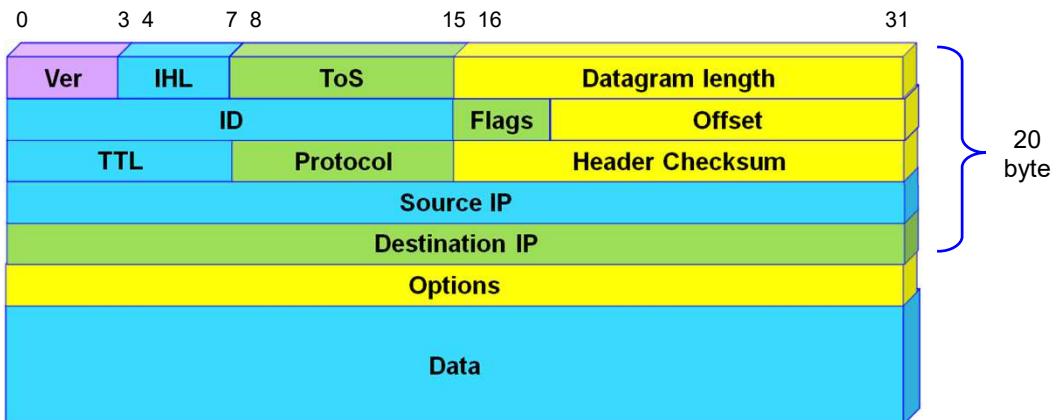
Assegnamento Indirizzi IP

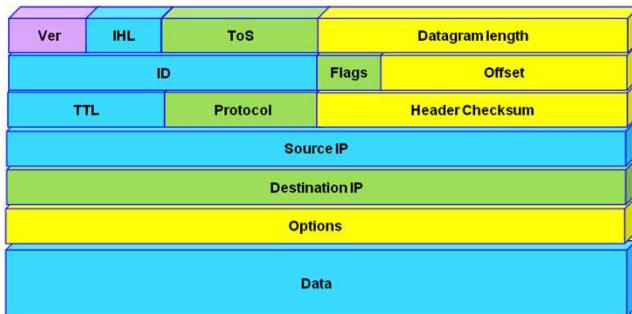
Ai RIR fanno capo i LIR (Local Internet Registries) per gestire gli indirizzi a livello locale.

GARR-LIR



Formato Datagram IPv4





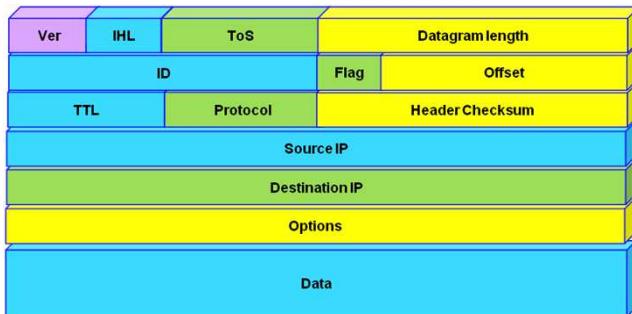
Version: 4 bit.

IHL: 4 bit - dimensione del preambolo, in gruppi da 4 byte

ToS: tipo di servizio: 3 bit di precedenza e 3 flag. Non è usato.

Datagram Length: lunghezza totale del datagram. 16 bit => 64kB

Formato Datagram IPv4

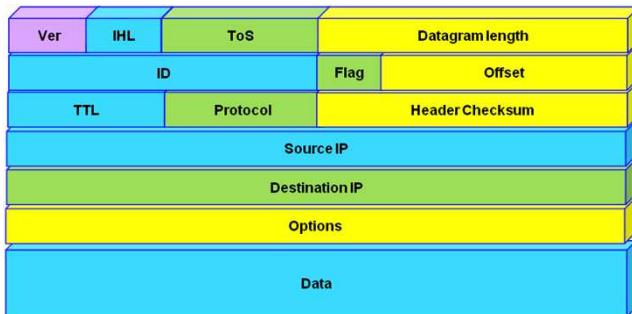


ID: identifica il datagram e i suoi frammenti.

FLAG (3 bit)

- **Not used**
- **DF** Don't fragment
- **MF**: More fragment

Fragment Offset: 13 bit => 8192 frammenti da almeno 8 byte



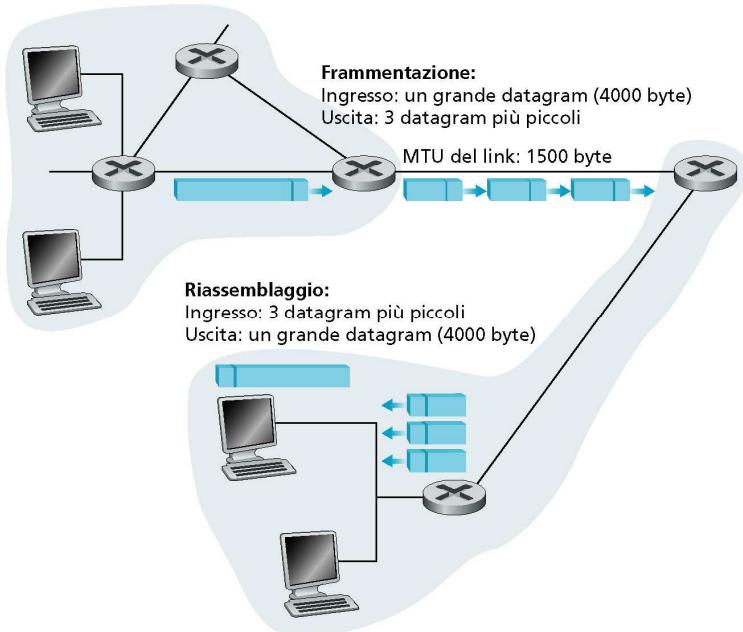
TTL: Time To Live

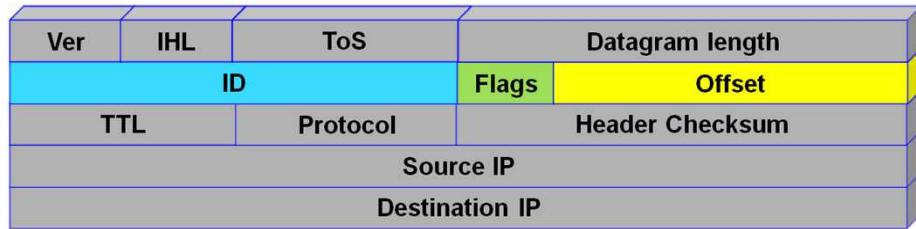
Protocol: specifica il protocollo di livello di trasporto

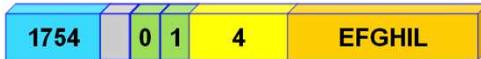
Header checksum: checksum del solo preambolo.

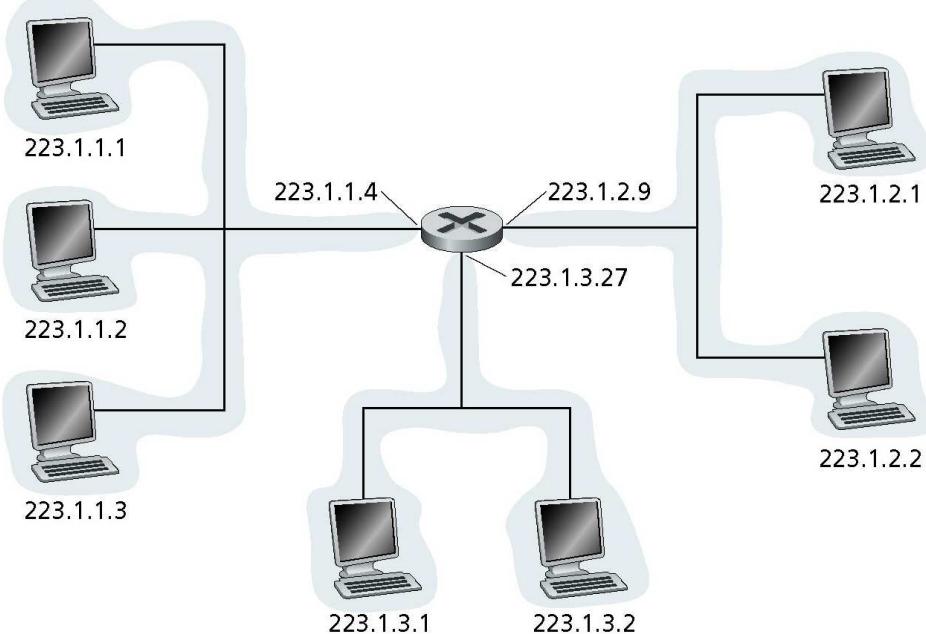
Options

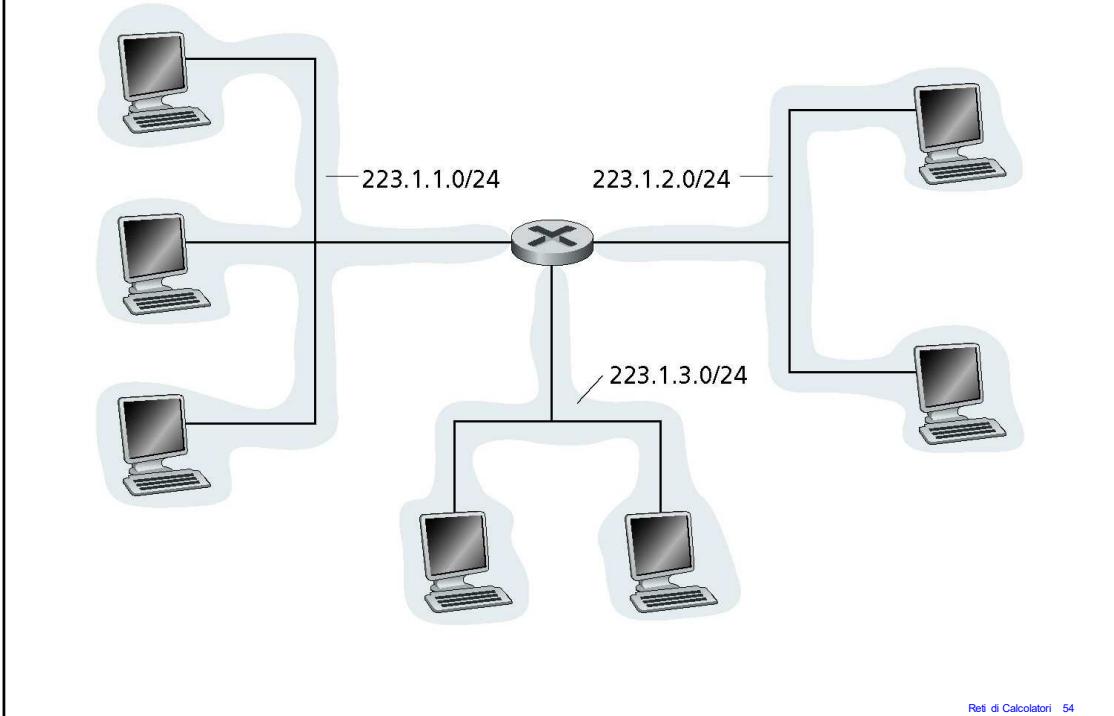
Frammentazione



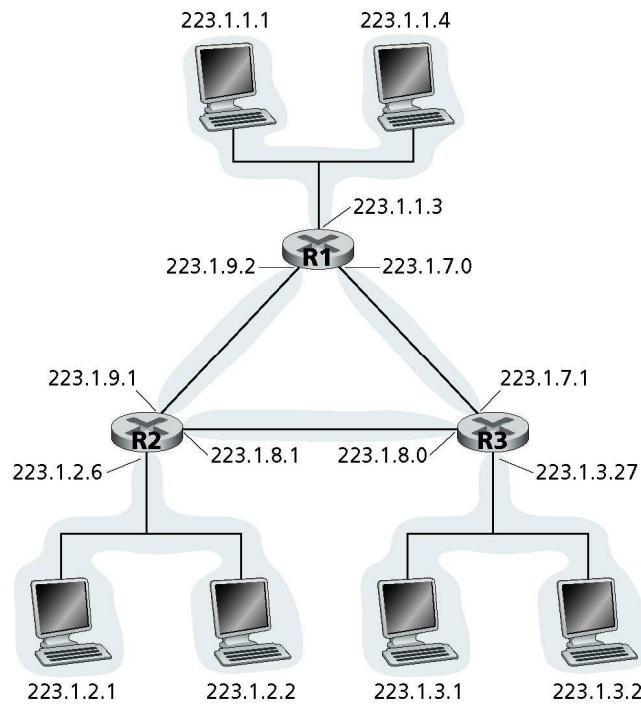




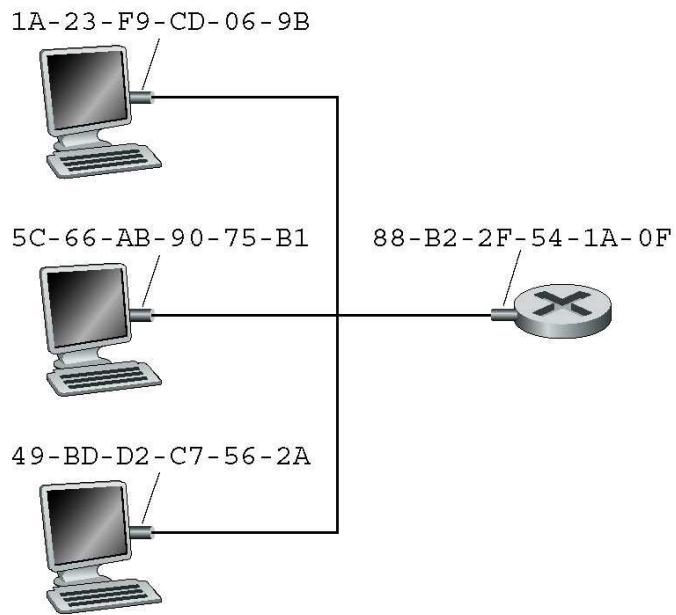


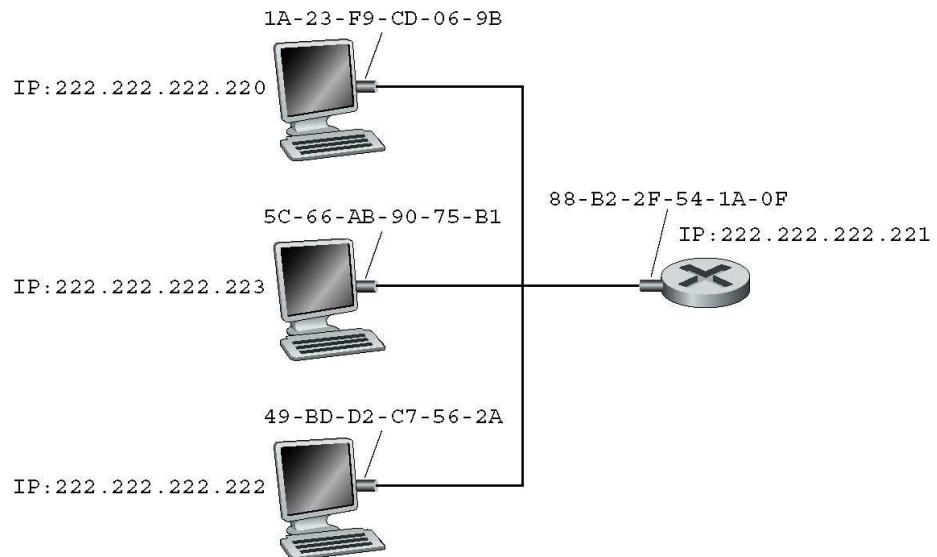


Indirizzamento IPv4

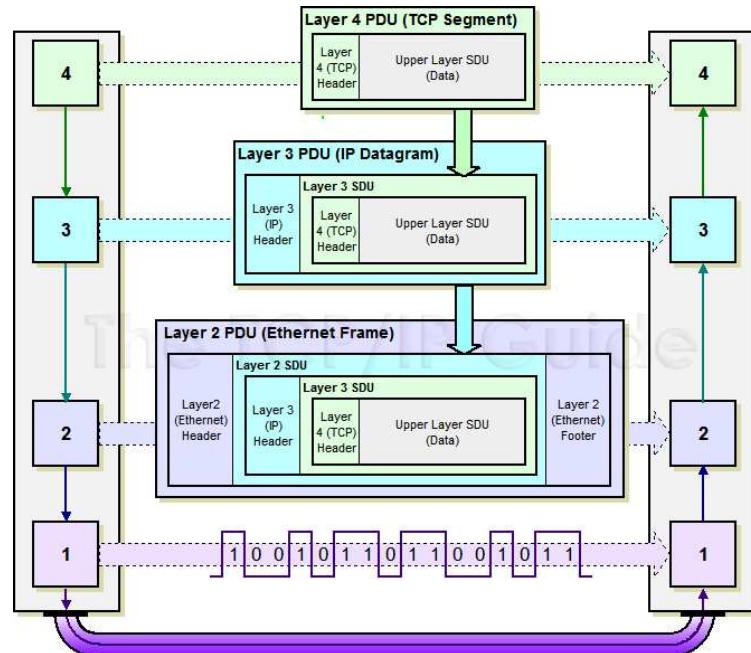


Reti di Calcolatori 55

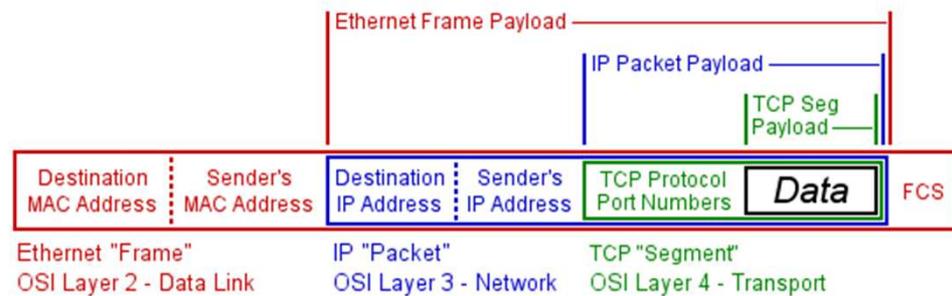




Indirizzi LAN e IPv4



Encapsulation Payloads



S.	Time	Source	Destination	Protocol	Length	Info
30	4.203963853	proxy03.mirror.garr.it	Toshiba-Z830.local	TCP	66	80 → 55492 [SYN, ACK]
31	4.203985642	Toshiba-Z830.local	proxy03.mirror.garr.it	TCP	54	55492 → 80 [ACK] Seq:
32	4.209664430	Toshiba-Z830.local	proxy03.mirror.garr.it	HTTP	802	GET / HTTP/1.1
33	4.213469694	mil04s25-in-f74.1e100.net	Toshiba-Z830.local	TCP	74	80 → 58828 [SYN, ACK]
34	4.213522703	Toshiba-Z830.local	mil04s25-in-f74.1e100.net	TCP	66	58828 → 80 [ACK] Seq:
35	4.222429261	proxy03.mirror.garr.it	Toshiba-Z830.local	TCP	60	80 → 55484 [ACK] Seq:

Frame 32: 802 bytes on wire (6416 bits), 802 bytes captured (6416 bits) on interface 0
Ethernet II, Src: Toshiba_dc:c7:5e (e8:e0:b7:c7:5e), Dst: Routerbo_13:e0:1c (4c:5e:0c:13:e0:1c)
Internet Protocol Version 4, Src: Toshiba-Z830.local (192.168.65.167), Dst: proxy03.mirror.garr.it (90.147.160.73)
Transmission Control Protocol, Src Port: 55484, Dst Port: 80, Seq: 1, Ack: 1, Len: 748

Hypertext Transfer Protocol

- ▶ GET / HTTP/1.1\r\n
- Host: mirror.garr.it\r\n
- User-Agent: Mozilla/5.0 (X11; Ubuntu; Linux x86_64; rv:65.0) Gecko/20100101 Firefox/65.0\r\n
- Accept: text/html,application/xhtml+xml,application/xml;q=0.9,image/webp,*/*;q=0.8\r\n
- Accept-Language: it-IT,it;q=0.8,en-US;q=0.5,en;q=0.3\r\n
- Accept-Encoding: gzip, deflate\r\n
- Referer: https://it.search.yahoo.com/\r\n
- Connection: keep-alive\r\n
- Cookie: __utma=225464966.2039547598.1552470368.1552470368.1552470368.1; __utmb=225464966.1.10.1552470368; __utmc=225
- Upgrade-Insecure-Requests: 1\r\n
- If-Modified-Since: Sat, 10 Nov 2018 08:45:07 GMT\r\n

```

000 4c 5e 0c 13 e0 1c e8 e0 b7 dc c7 5e 08 00 45 00  LA..... .^.E.
010 03 14 93 d1 40 00 40 06 a6 e6 c8 a8 41 a7 5a 93  ...@. A Z-
020 a0 49 d8 bc 00 50 60 16 5a be ca 22 4d 02 50 18  I - P . Z .."M-P-
030 00 e5 00 33 00 90 47 45 54 29 2f 20 48 54 54 50  ..3..GE T / HTTP
040 2f 31 2e 31 0d 0a 48 6f 73 74 3a 20 6d 69 72 72 /1.1.Ho st: mirr
050 6f 72 2e 67 61 72 72 2e 69 74 0d 0a 55 73 65 72 or.garr.it>User
060 2d 41 67 65 6e 74 3a 20 4d 6f 7a 69 6c 6c 61 2f -Agent: Mozilla/
070 35 2e 30 20 28 58 31 31 3b 20 55 62 75 6e 74 75 5.0 (X11; Ubuntu
080 3b 20 4c 69 6e 75 78 20 78 38 36 5f 36 34 3b 20 ; Linux x86_64;
090 72 76 3a 3b 35 2e 30 29 20 47 65 63 6b 6f 2f 32 rv:65.0) Gecko/2
0a0 30 31 30 30 31 30 31 29 46 69 72 65 66 6f 78 2f 0100101 Firefox/
0b0 36 35 2e 30 0d 0a 41 63 63 65 70 74 3a 20 74 65 65.0.Ac cept: te
0c0 78 74 2f 68 74 6d 6c 2c 61 70 70 6c 69 63 61 74 xt/html, applicat
0d0 69 6f 6e 2f 78 68 74 6d 6c 2b 78 6d 6c 2c 61 70 ion/xhtml+xml,ap
0e0 70 6c 69 63 61 74 69 6f 6e 2f 78 6d 6c 3b 71 3d plicatio n/xml;q=
0f0 30 2e 39 2c 69 6d 61 67 65 2f 77 65 62 70 2c 2a 0.9,imag e/webp,*
100 2f 2a 3b 71 3d 30 2e 38 0d 0a 41 63 63 65 70 74 /*;q=0.8 .Accept

```

Indirizzi LAN e IPv4

S.	Time	Source	Destination	Protocol	Length	Info
30	4.203963853	proxy03.mirror.garr.it	Toshiba-Z830.local	TCP	66	80 → 55492 [SYN, ACK]
31	4.203985642	Toshiba-Z830.local	proxy03.mirror.garr.it	TCP	54	55492 → 80 [ACK] Seq:
32	4.209664430	Toshiba-Z830.local	proxy03.mirror.garr.it	HTTP	802	GET / HTTP/1.1
33	4.213469694	mil04s25-in-f74.1e100.net	Toshiba-Z830.local	TCP	74	80 → 58828 [SYN, ACK]
34	4.213522703	Toshiba-Z830.local	mil04s25-in-f74.1e100.net	TCP	66	58828 → 80 [ACK] Seq:
35	4.222429261	proxy03.mirror.garr.it	Toshiba-Z830.local	TCP	66	80 → 55484 [ACK] Seq:

Frame 32: 802 bytes on wire (6416 bits), 802 bytes captured (6416 bits) on interface a
Ethernet II, Src: Toshiba_dc:c7:5e (e8:e0:b7:dc:c7:5e), Dst: Routerbo_13:e0:1c (4c:5e:0c:13:e0:1c)
Internet Protocol Version 4, Src: Toshiba-Z830.local (192.168.65.167), Dst: proxy03.mirror.garr.it (90.147.160.73)
Transmission Control Protocol, Src Port: 55484, Dst Port: 80, Seq: 1, Ack: 1, Len: 748

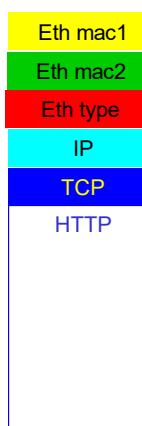
Hypertext Transfer Protocol

▶ GET / HTTP/1.1\r\n

Host: mirror.garr.it\r\nUser-Agent: Mozilla/5.0 (X11; Ubuntu; Linux x86_64; rv:65.0) Gecko/20100101 Firefox/65.0\r\nAccept: text/html,application/xhtml+xml,application/xml;q=0.9,image/webp,*/*;q=0.8\r\nAccept-Language: it-IT,it;q=0.8,en-US;q=0.5,en;q=0.3\r\nAccept-Encoding: gzip, deflate\r\nReferer: https://it.search.yahoo.com/\r\nConnection: keep-alive\r\n

000 4c 5e 9c 13 e0 1c e8 e0 b7 de c7 5e 08 00 45 00
010 03 14 93 d1 40 00 40 06 a6 e6 c0 a8 41 a7 5a 93
020 a0 49 d8 bc 00 50 60 10 5a be ca 22 4d 02 50 18
030 00 e5 00 33 00 00 47 45 54 20 2f 20 48 54 54 50
040 2f 31 2e 31 0d 0a 48 6f 73 74 3a 20 6d 69 72 72
050 6f 72 2e 67 61 72 72 2e 69 74 9d 0a 55 73 65 72
060 2d 41 67 65 6e 74 3a 20 4d 6f 7a 69 6c 6c 61 2f
070 35 2e 30 20 28 58 31 31 3b 20 55 62 75 6e 74 75
080 3b 20 4c 69 6e 75 78 20 78 38 36 5f 36 34 3b 20
090 72 76 3a 36 35 2e 39 29 20 47 65 63 6b 6f 2f 32
0a0 30 31 30 30 31 30 31 20 46 69 72 65 66 6f 78 2f
0b0 36 35 2e 30 0d 0a 41 63 63 65 70 74 3a 20 74 65
0c0 78 74 2f 68 74 6d 6c 2c 61 70 70 6c 69 63 61 74
0d0 69 6f 6e 2f 78 68 74 6d 6c 2b 78 6d 6c 2c 61 70
0e0 70 6c 69 63 61 74 69 6f 6e 2f 78 6d 6c 3b 71 3d
0f0 39 2e 39 2c 69 6d 61 67 65 2f 77 65 62 70 2c 2a
100 2f 2a 3b 71 3d 30 2e 38 0d 0a 41 63 63 65 70 74

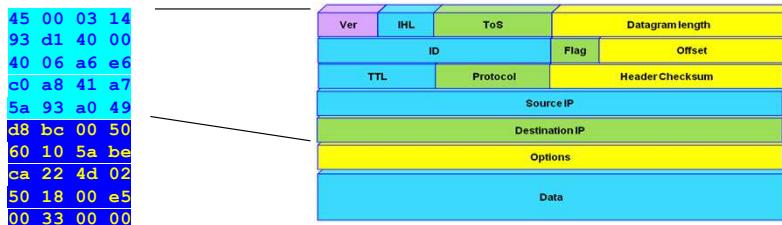
Reti di Calcolatori 61



Indirizzi LAN e IPv4

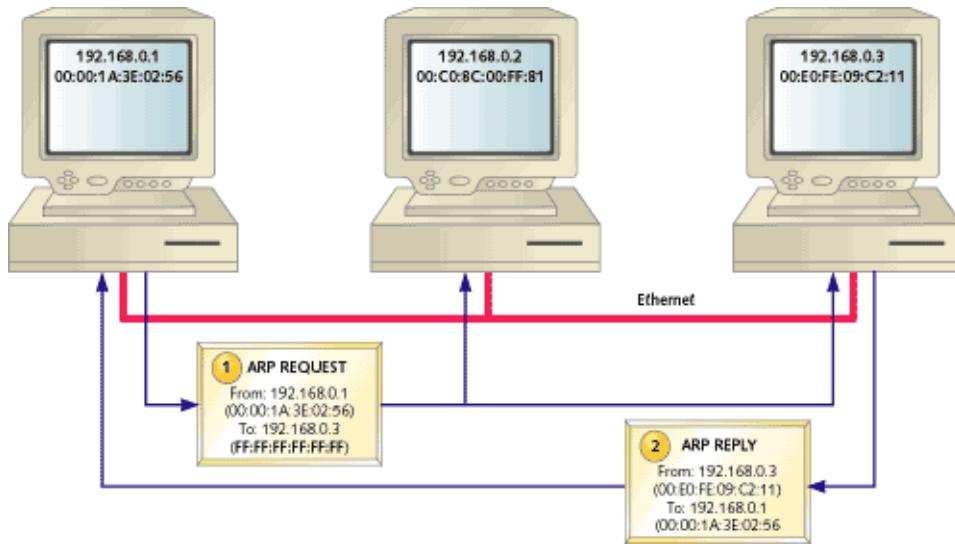
No.	Time	Source	Destination	Protocol	Length	Info
30	4.203963853	proxy03.mirror.garr.it	Toshiba-Z830.local	TCP	66	80 → 55492 [SYN, ACK]
31	4.203985642	Toshiba-Z830.local	proxy03.mirror.garr.it	TCP	54	55492 → 80 [ACK] Seq:
32	4.203964430	Toshiba-Z830.local	proxy03.mirror.garr.it	HTTP	802	GET / HTTP/1.1
33	4.213469694	mil04s25-in-f74.ie100.net	Toshiba-Z830.local	TCP	74	80 → 58828 [SYN, ACK]
34	4.213522763	Toshiba-Z830.local	mil04s25-in-f74.ie100.net	TCP	66	58828 → 80 [ACK] Seq:
35	4.222429261	proxy03.mirror.garr.it	Toshiba-Z830.local	TCP	60	80 → 55484 [ACK] Seq:

Frame 32: 802 bytes on wire (6416 bits), 802 bytes captured (6416 bits) on interface 0
 Ethernet II, Src: Toshiba_dc:c7:5e (e8:e0:b7:dc:c7:5e), Dst: Routerbo_13:e0:1c (4c:5e:0c:13:e0:1c)
 Internet Protocol Version 4, Src: Toshiba-Z830.local (192.168.65.167), Dst: proxy03.mirror.garr.it (90.147.160.73)
 Transmission Control Protocol, Src Port: 55484, Dst Port: 80, Seq: 1, Ack: 1, Len: 748



Ver 04 IHL 05 ToS 03 DL 0314 (788 = 802-6-6-2)
ID 93D1 Flag + Offset 010 0000000000000000
TTL 40 (64) Protocol 06 (TCP) HC A6E6
Source IP C0 (192) A8 (168) 41 (065) A7 (167)
Dest IP 5A (090) 93 (147) A0 (160) 49 (073)

ARP request

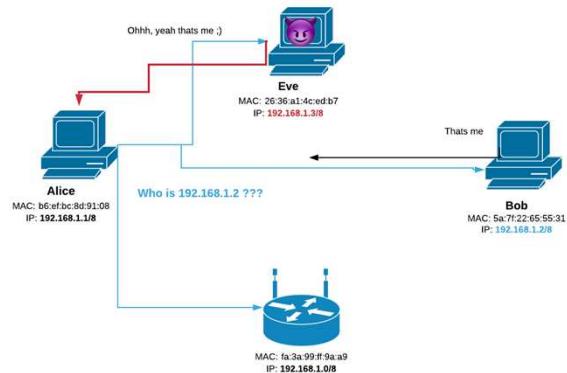


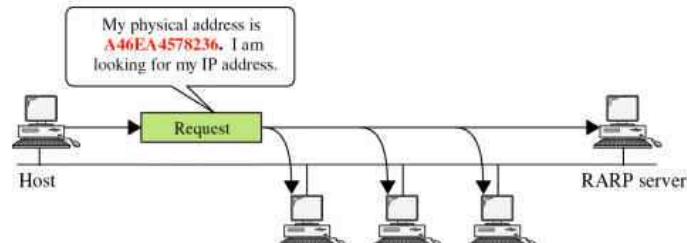
Indirizzo IP	Indirizzo LAN	TTL
222.222.222.221	88-B2-2F-54-1A-0F	13:45:00
222.222.222.223	5C-66-AB-90-75-B1	13:52:00

- ARP non prevede autenticazione.
- ARP è un protocollo senza stato: una ARP Reply può essere mandata senza una precedente ARP Request.
- Un host che riceve una pacchetto ARP deve aggiornare la sua ARP cache.

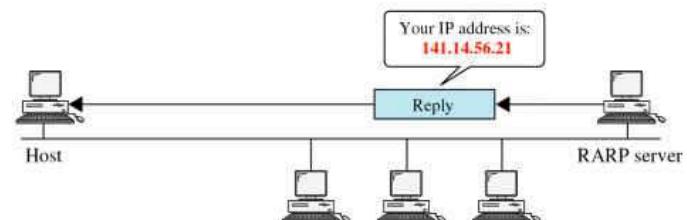
Conseguenze

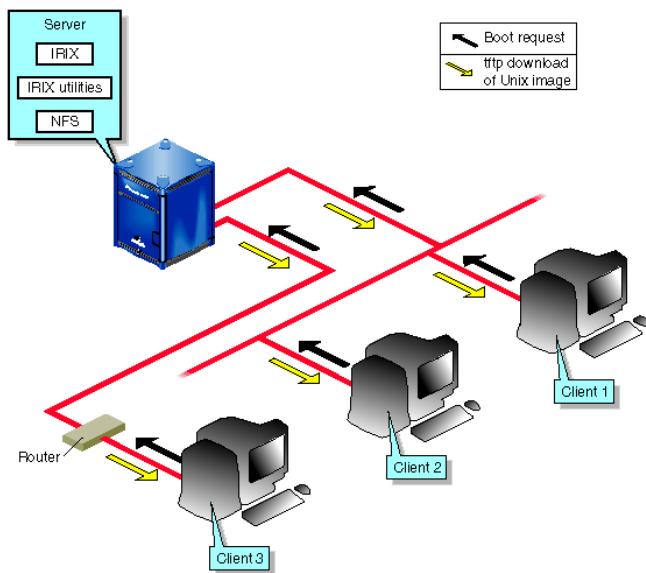
È possibile dirottare il traffico IP.

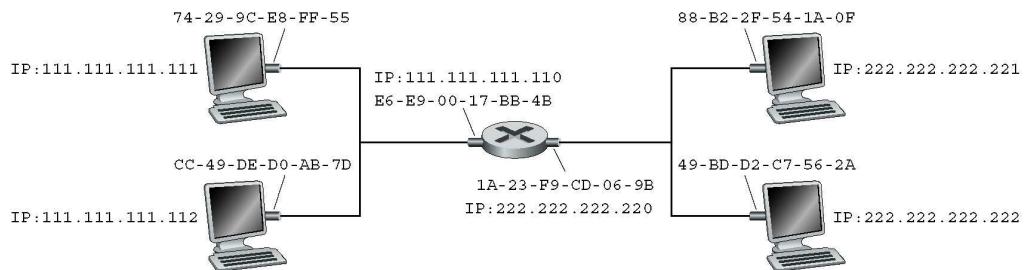




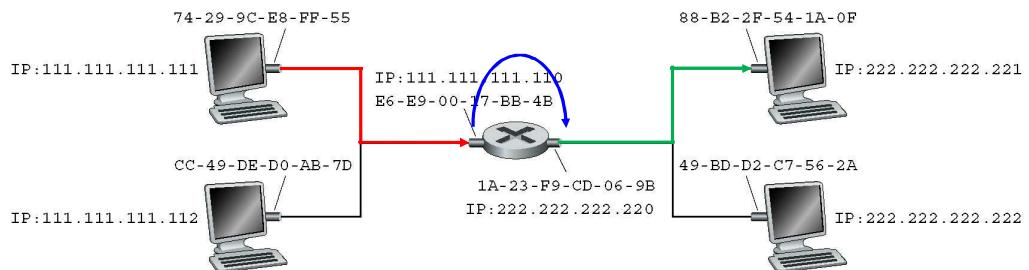
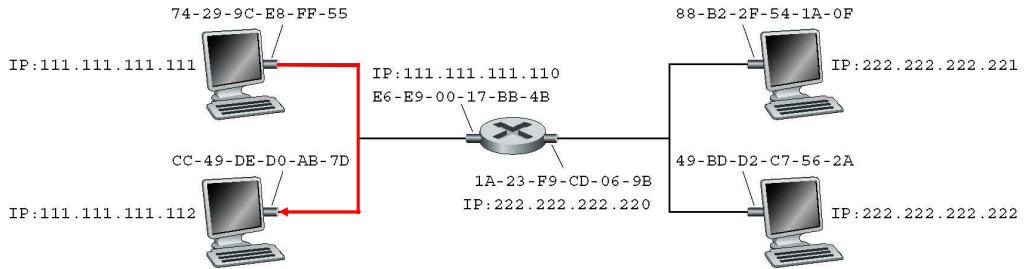
a. RARP request is broadcast







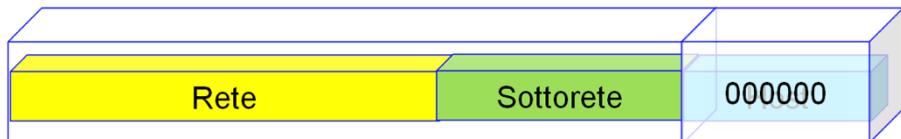
LAN interconnesse



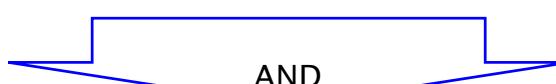
CIDR (Classless Inter-Domain Routing)

VLSM (Variable Length Subnet Masking)

a.b.c.d/x



10010111.01100001.11111100.01000010 151.097.252.066
11111111.11111111.11111110.00000000 255.255.254.000

 AND
10010111.01100001.11111100.00000000 151.097.252.000



151.097.252.000/23

10010111.01100001.11111100.00000000

10010111.01100001.11111100.00000001 151.097.252.001

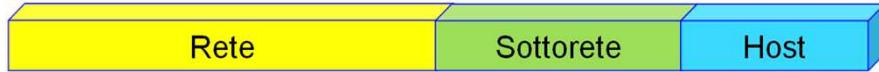
10010111.01100001.11111100.00000010 151.097.252.002

...

10010111.01100001.11111101.11111110 151.097.253.254

151.97.252.0/23 Tutta la sottorete

151.97.253.255 Broadcast nella sottorete



191.15.12.0/30

10111111.00001111.00001100.000000**00** **network**

10111111.00001111.00001100.000000**01**

10111111.00001111.00001100.000000**10**

10111111.00001111.00001100.000000**11** **broadcast**

191.15.12.24/30

10111111.00001111.00001100.0001**1000** **network**

191.15.12.25

191.15.12.26

191.15.12.27 10111111.00001111.00001100.00011011****



/31

A 31-bit subnet mask is often used for an interface that is the endpoint of a point-to-point network. The use of 31-bit subnet masks for IPv4 point-to-point links is described in RFC 3021.

/32

A 32-bit subnet mask defines a network with only one IP address.

Netmask

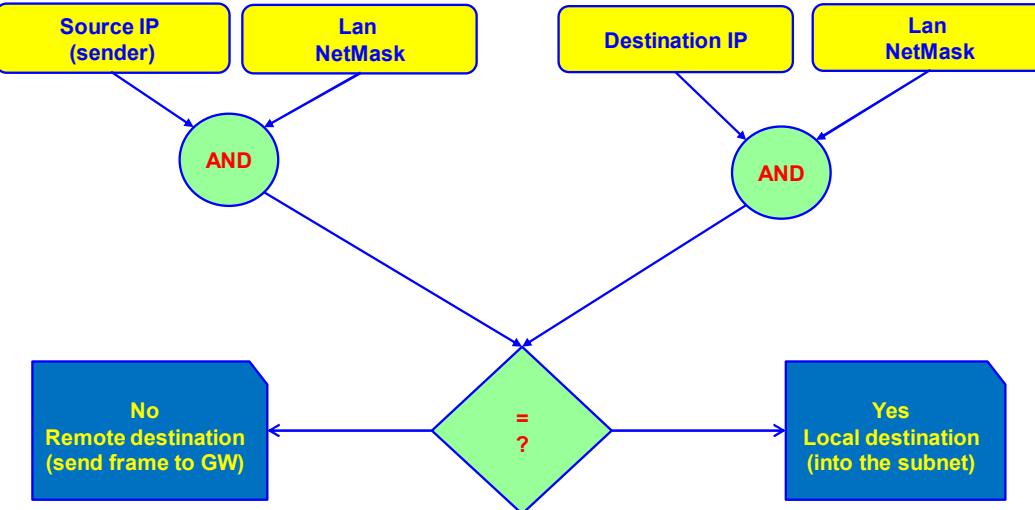


Tabella host

Tabella di rilancio in A		
Rete di destinazione	Router successivo	Numero salti
223.1.1.0/24		1
223.1.2.0/24	223.1.1.4	2
223.1.3.0/24	223.1.1.4	2

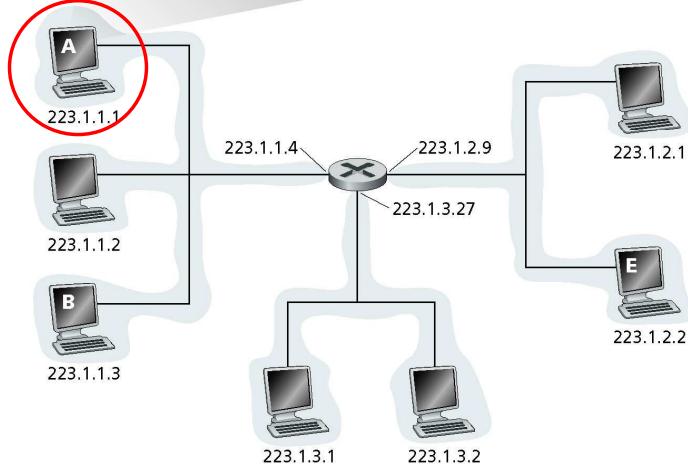
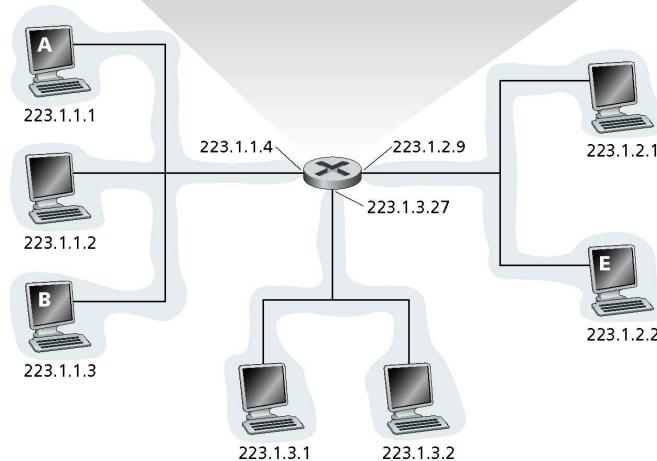


Tabella di routing

Tabella di rilancio nel router			
Rete di destinaz.	Router success.	N. hop	Interfaccia
223.1.1.0/24	—	1	223.1.1.4
223.1.2.0/24	—	1	223.1.2.9
223.1.3.0/24	—	1	223.1.3.27



Example: routing table

Scheda Ethernet Ethernet:

```
Suffisso DNS specifico per connessione: dmi.unict.it
Indirizzo IPv6 locale rispetto al collegamento : fe80::2430:ee60:719:1e6c%6
Indirizzo IPv4. . . . . : 192.168.65.156
Subnet mask . . . . . : 255.255.255.0
Gateway predefinito . . . . . : 192.168.65.1
```

Scheda Ethernet Ethernet 3:

```
Suffisso DNS specifico per connessione:
Indirizzo IPv6 locale rispetto al collegamento : fe80::c5f:4e6:d1c1:c3d2%39
Indirizzo IPv4. . . . . : 192.168.56.1
Subnet mask . . . . . : 255.255.255.0
Gateway predefinito . . . . . :
```

Scheda LAN wireless Wi-Fi:

```
Suffisso DNS specifico per connessione: wluct.unict.it
Indirizzo IPv6 locale rispetto al collegamento : fe80::b176:bbcd:561:97ae%54
Indirizzo IPv4. . . . . : 151.97.62.134
Subnet mask . . . . . : 255.255.248.0
Gateway predefinito . . . . . : 151.97.56.1
```

Stato supporto. : Supporto disconnesso
Suffisso DNS specifico per connessione: wluct.unict.it

Example: routing table

```
C:\> route print
=====
Elenco interfacce
55.....Wintun Userspace Tunnel
 6...00 1c c2 4b 8a b3 .....Realtek USB GbE Family Controller
 39...0a 00 27 00 00 27 .....VirtualBox Host-Only Ethernet Adapter
 18...00 ff 51 02 c2 ab .....TAP-Windows Adapter V9
 35...38 fc 98 b0 b3 56 .....Microsoft Wi-Fi Direct Virtual Adapter
  4...3a fc 98 b0 b3 55 .....Microsoft Wi-Fi Direct Virtual Adapter #2
 54...38 fc 98 b0 b3 55 .....Intel(R) Wi-Fi 6 AX200 160MHz
 17...38 fc 98 b0 b3 59 .....Bluetooth Device (Personal Area Network)
  1.....Software Loopback Interface 1
 28...94 8f 85 51 5b 53 .....Generic Mobile Broadband Adapter
=====
```

Example: routing table

IPv4 Tabella route

Route attive:

Indirizzo rete	Mask	Gateway	Interfaccia	
Metrica				
0.0.0.0	0.0.0.0	151.97.56.1	151.97.62.134	45
127.0.0.0	255.0.0.0	On-link	127.0.0.1	331
127.0.0.1	255.255.255.255	On-link	127.0.0.1	331
127.255.255.255	255.255.255.255	On-link	127.0.0.1	331
151.97.56.0	255.255.248.0	On-link	151.97.62.134	301
151.97.62.134	255.255.255.255	On-link	151.97.62.134	301
151.97.63.255	255.255.255.255	On-link	151.97.62.134	301
192.168.56.0	255.255.255.0	On-link	192.168.56.1	281
192.168.56.1	255.255.255.255	On-link	192.168.56.1	281
192.168.56.255	255.255.255.255	On-link	192.168.56.1	281
224.0.0.0	240.0.0.0	On-link	127.0.0.1	331
224.0.0.0	240.0.0.0	On-link	192.168.56.1	281
224.0.0.0	240.0.0.0	On-link	151.97.62.134	301
255.255.255.255	255.255.255.255	On-link	127.0.0.1	331
255.255.255.255	255.255.255.255	On-link	192.168.56.1	281
255.255.255.255	255.255.255.255	On-link	151.97.62.134	301

Route permanenti:

Nessuna

Example: routing table

IPv4 Tabella route
Route attive:

Indirizzo rete	Mask	Gateway	Interfaccia	
Metrica				
0.0.0.0	0.0.0.0	192.168.65.1	192.168.65.156	25
127.0.0.0	255.0.0.0	On-link	127.0.0.1	331
127.0.0.1	255.255.255.255	On-link	127.0.0.1	331
127.255.255.255	255.255.255.255	On-link	127.0.0.1	331
192.168.56.0	255.255.255.0	On-link	192.168.56.1	281
192.168.56.1	255.255.255.255	On-link	192.168.56.1	281
192.168.56.255	255.255.255.255	On-link	192.168.56.1	281
192.168.65.0	255.255.255.0	On-link	192.168.65.156	281
192.168.65.156	255.255.255.255	On-link	192.168.65.156	281
192.168.65.255	255.255.255.255	On-link	192.168.65.156	281
224.0.0.0	240.0.0.0	On-link	127.0.0.1	331
224.0.0.0	240.0.0.0	On-link	192.168.56.1	281
224.0.0.0	240.0.0.0	On-link	192.168.65.156	281
255.255.255.255	255.255.255.255	On-link	127.0.0.1	331
255.255.255.255	255.255.255.255	On-link	192.168.56.1	281
255.255.255.255	255.255.255.255	On-link	192.168.65.156	281

Route permanenti:

Nessuna

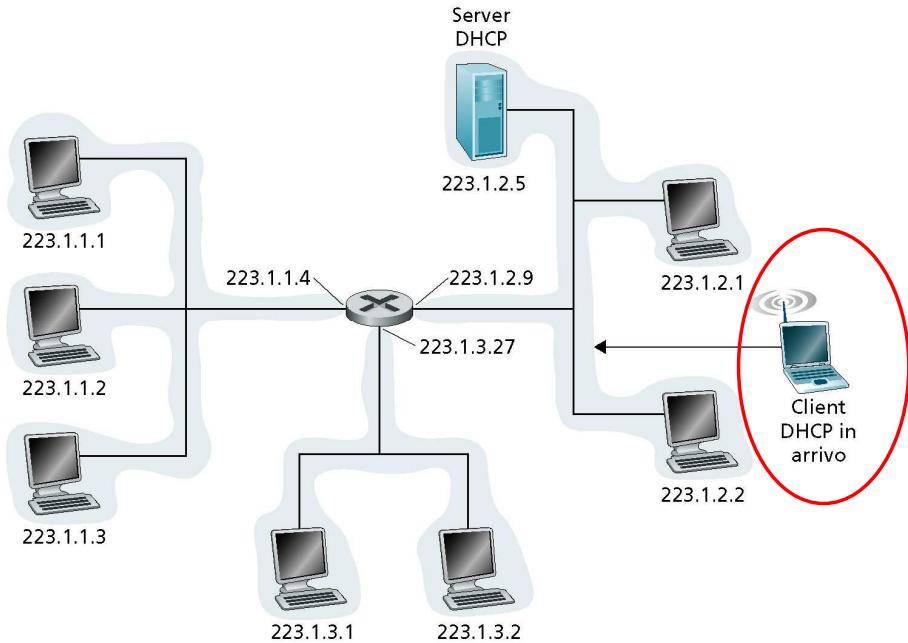
Example: routing table

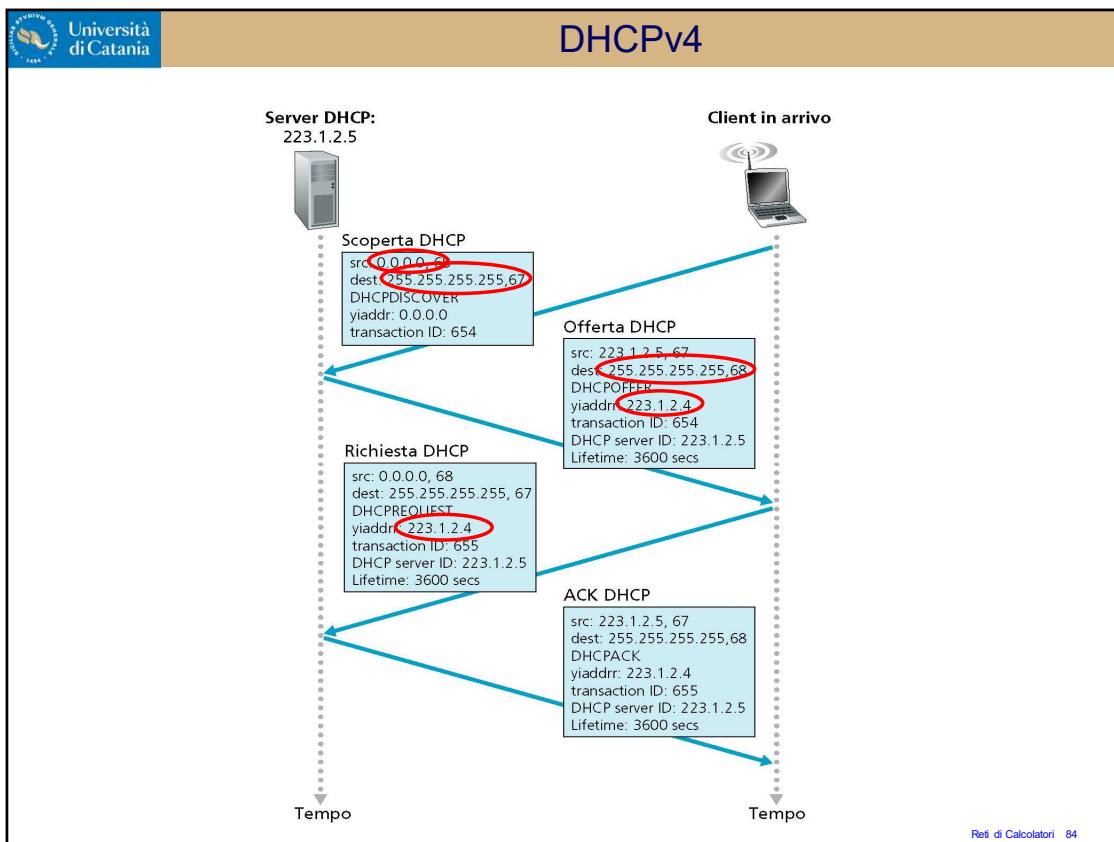
IPv6 Tabella route**Route attive:**

Interf	Metrica	Rete	Destinazione	Gateway
1	331	::1/128		On-link
39	281	fe80::/64		On-link
6	281	fe80::/64		On-link
39	281	fe80::c5f:4e6:d1c1:c3d2/128		On-link
6	281	fe80::2430:ee60:719:1e6c/128		On-link
1	331	ff00::/8		On-link
39	281	ff00::/8		On-link
6	281	ff00::/8		On-link

Route permanenti:**Nessuna**

Assegnamento degli indirizzi: DHCPv4

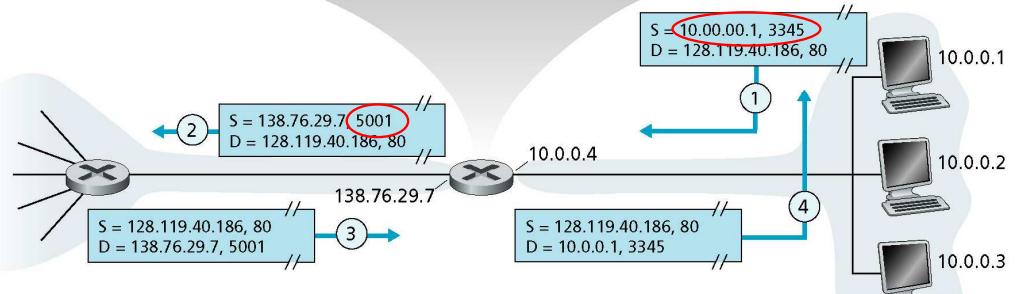




Offerta DHCP

```
src: 223.1.2.5, 67
dest: 255.255.255.255,68
DHCPOFFER
yiaddr: 223.1.2.4
transaction ID: 654
DHCP server ID: 223.1.2.5
Lifetime: 3600 secs
```

Tabella di traduzione NAT	
Lato WAN	Lato LAN
138.76.29.7, 5001	10.00.00.1, 3345
...	...



“If Y2K was technology's ticking time bomb,
then IPv6 is a slow, deadly gas leak.”

Limiti di IPv4

- Esaurimento dello spazio degli indirizzi
- Scalabilità del routing
- Nuovi servizi

“I think there is a world market for maybe five computers”

– Thomas Watson (IBM), 1943

“640K ought to be enough for anybody”

– Bill Gates (MS), 1981

“32 bits should be enough address space for the Internet”

– Vint Cerf (ARPA), 1977

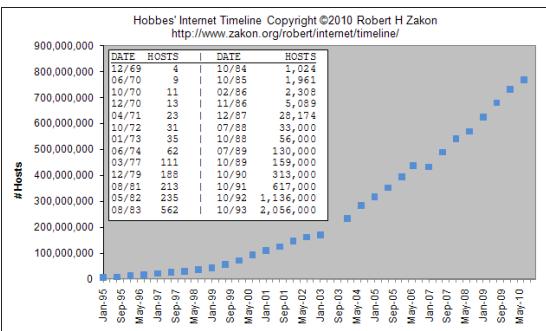
IPv4 Address Allocation

IPv4 addresses Available in October 17, 2010 according to the RIPE															
0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
32	33	34	35	36	37	38	39	40	41	42	43	44	45	46	47
48	49	50	51	52	53	54	55	56	57	58	59	60	61	62	63
64	65	66	67	68	69	70	71	72	73	74	75	76	77	78	79
80	81	82	83	84	85	86	87	88	89	90	91	92	93	94	95
96	97	98	99	100	101	102	103	104	105	106	107	108	109	110	111
112	113	114	115	116	117	118	119	120	121	122	123	124	125	126	127
128	129	130	131	132	133	134	135	136	137	138	139	140	141	142	143
144	145	146	147	148	149	150	151	152	153	154	155	156	157	158	159
160	161	162	163	164	165	166	167	168	169	170	171	172	173	174	175
176	177	178	179	180	181	182	183	184	185	186	187	188	189	190	191
192	193	194	195	196	197	198	199	200	201	202	203	204	205	206	207
208	209	210	211	212	213	214	215	216	217	218	219	220	221	222	223
224	225	226	227	228	229	230	231	232	233	234	235	236	237	238	239
240	241	242	243	244	245	246	247	248	249	250	251	252	253	254	255

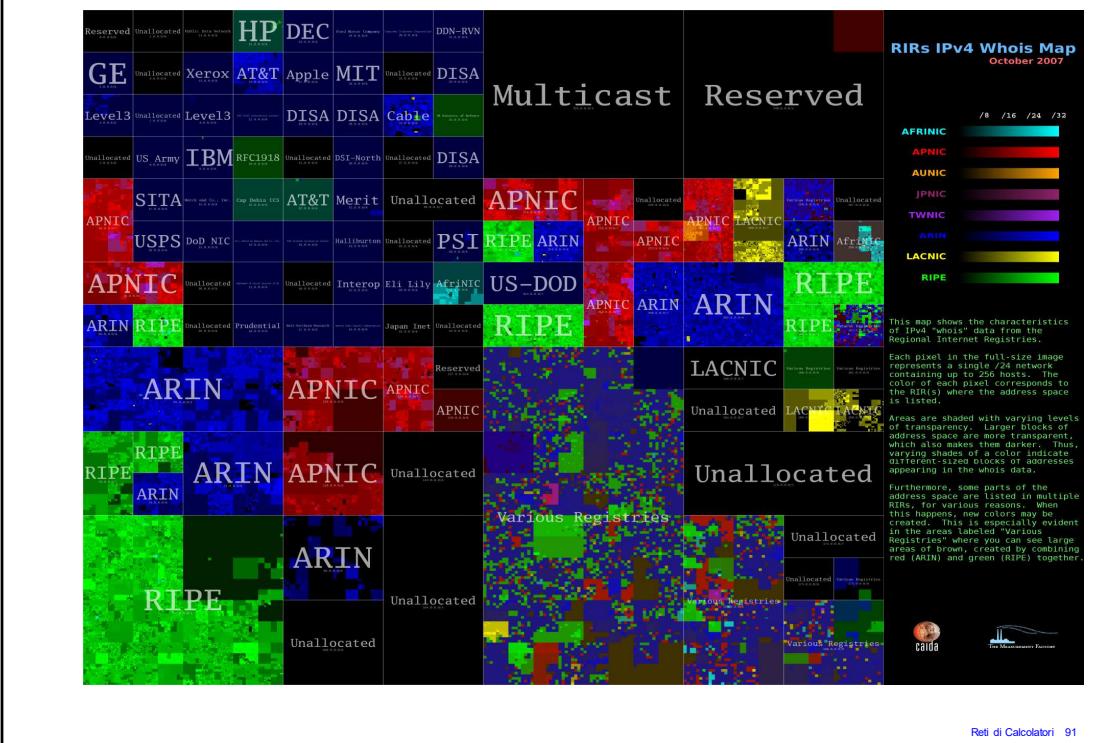
xx Used
xx Available
xx Notuseable

xx Allocated in Jan. 2010
xx Allocated in May 2010
xx Allocated in Oct. 2010
xx Allocated in Feb. 2010
xx Allocated in June 2010
xx Allocated in Apr. 2010
xx Allocated in Aug. 2010

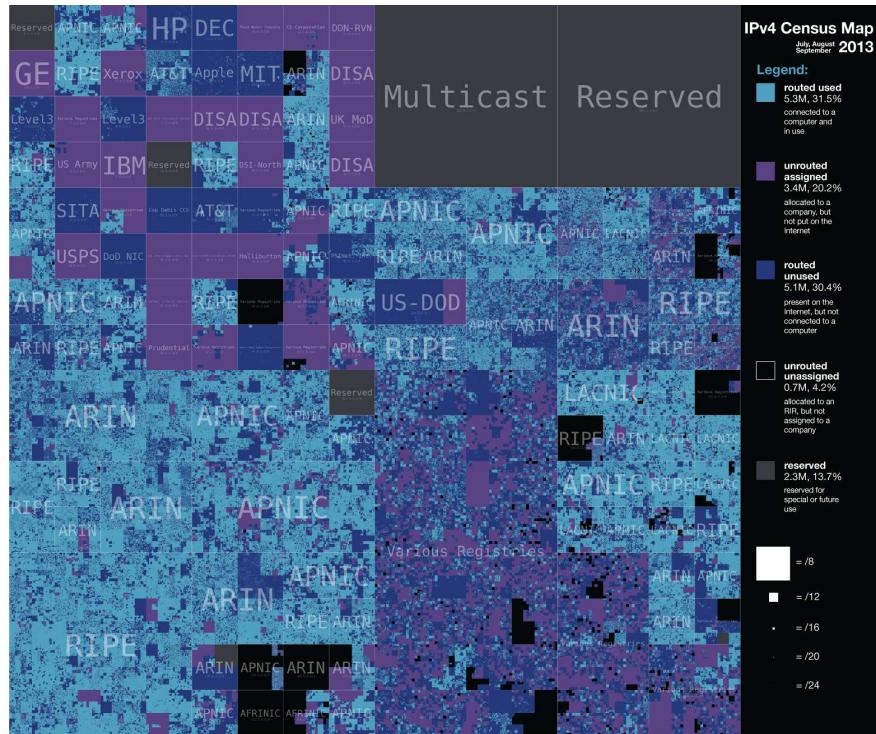
Total 256
Free 12
Percent 4.69%



IPv4 Address Allocation

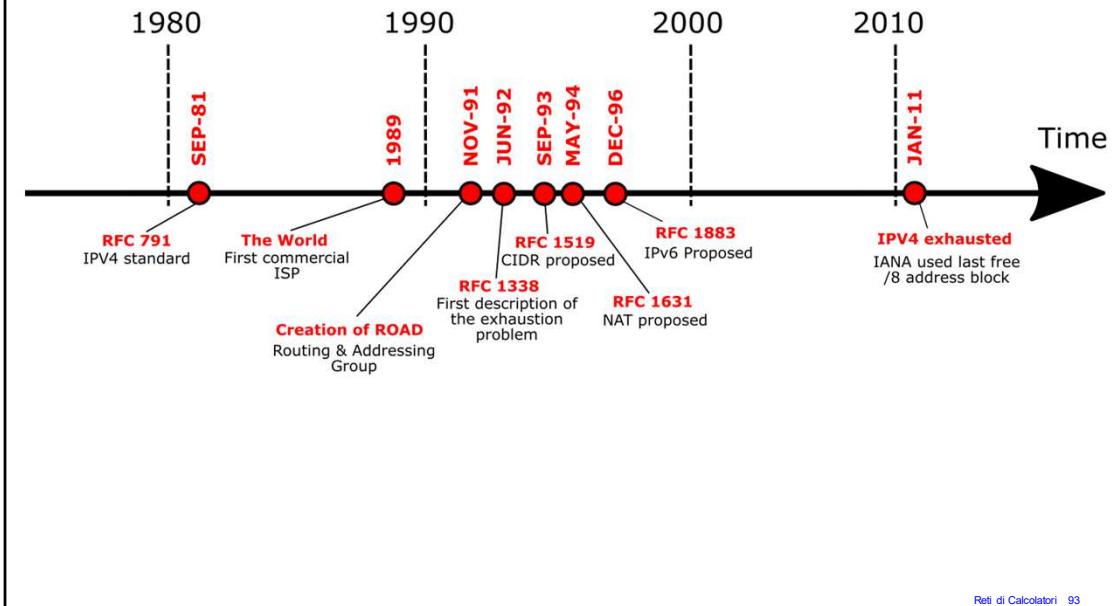


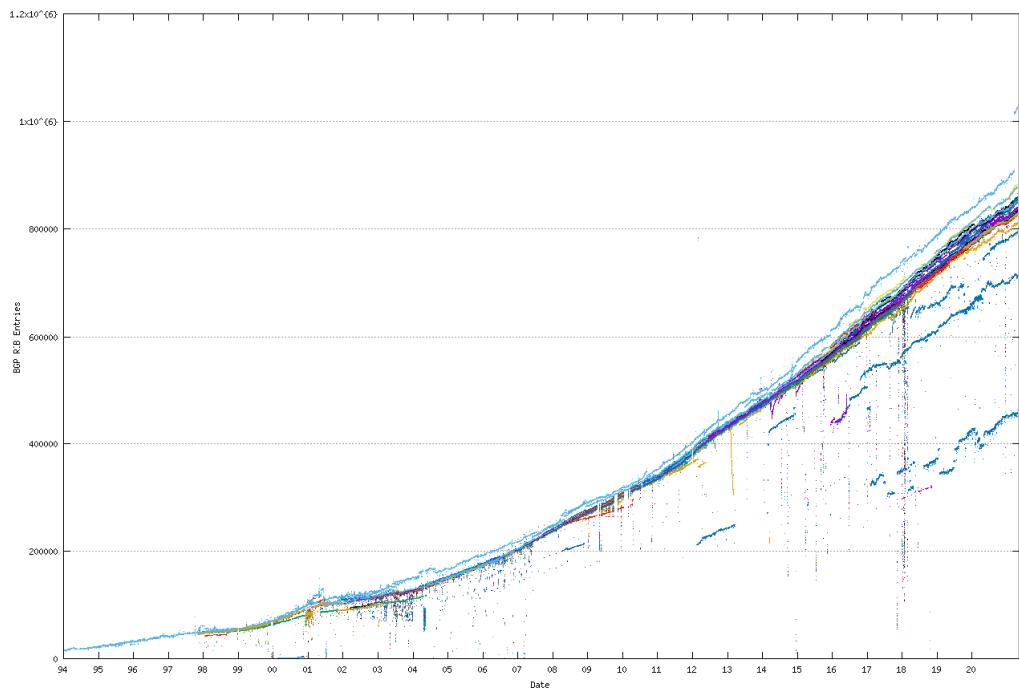
IPv4 Address Allocation



Reti di Calcolatori 92

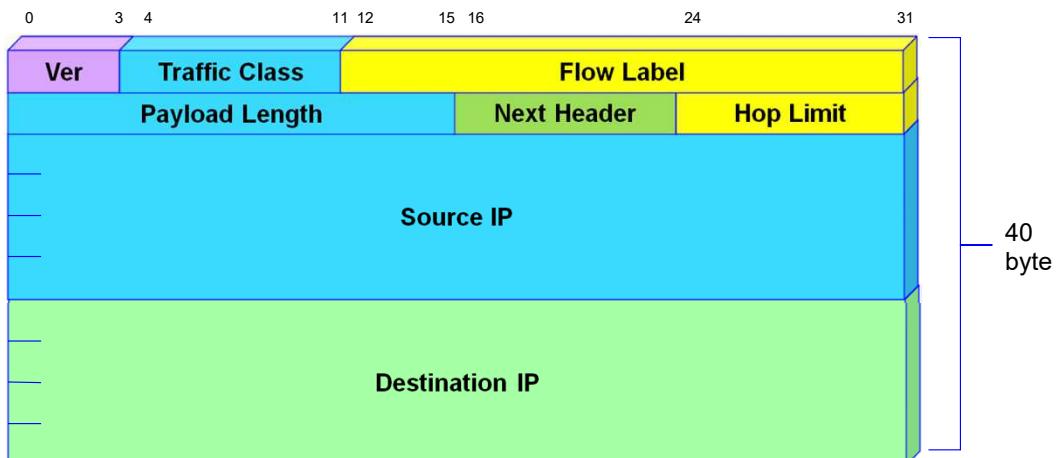
IPv4 Address Allocation





Nuovi Servizi richiesti :

- Sicurezza
- Autoconfigurazione (Plug & Play)
- Gestione della Qualità del Servizio (QoS)
- Indirizzamento Multicast
- Indirizzamento host mobili

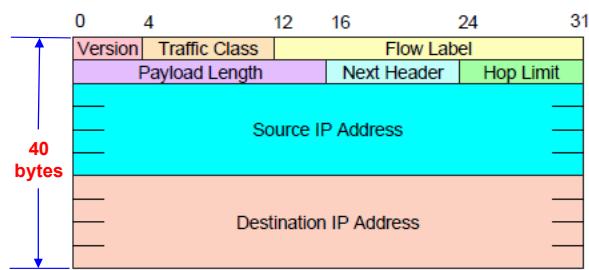


$$2^{128} = 340282366920938463463374607431768211456 \approx 3 \times 10^{38}$$

$$2^{32} = 4294967296 \approx 4 \times 10^9$$

Sup. terrestre $\approx 5.1 \times 10^{14} \text{ m}^2 \Rightarrow 6.6 \times 10^{23} \text{ indirizzi IPv6 per m}^2$
 Stelle nell'universo (osservabile) $\approx 10^{24} \Rightarrow 10^{14} \text{ indirizzi per stella}$

IPv6



Rimossi

- ID, Flags, Offset
- ToS, HLEN
- Header Checksum

Cambiati

- Total length \Rightarrow payload
- Protocol \Rightarrow next header
- TTL \Rightarrow hop limit

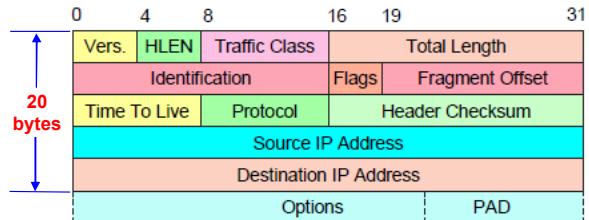
Aggiunti

- Traffic class
- Flow label

Espansi

- address: 32 \Rightarrow 128 bits

IPv4



0000 0000	Riservati (includono IPv4)
0000 001	Indirizzi OSI NSAP (deprecated)
0000 010	Indirizzi Novell IPX
001	Indirizzi unicast globali
010	Indirizzi per service providers
100	Indirizzi geografici
1111 1110 10	Indirizzi unicast link-local
1111 1110 11	Indirizzi unicast site-link (deprecated)
1111 1111	Multicast

8000:0000:0000:0000:0123:4567:89AB:CDEF

8000::123:4567:89AB:CDEF

2001:760:4804:0:0:0:0:105

2001:760:4804::/48

::151.97.1.1 indirizzo IPv4 mappato su IPv6

::1 loopback

:: non specificato

Global Unicast addresses

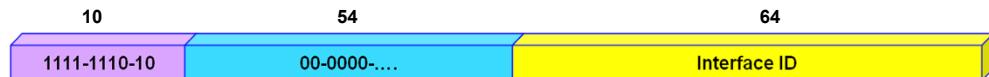


2 001:760:4804: 0 : 0 :0:0:105

Global Routing Prefix: identifica il sito.

Subnet ID: viene gestito all'interno del sito

Indirizzi locali: **FE80 :: /64** (1111-1110-1000-0000)



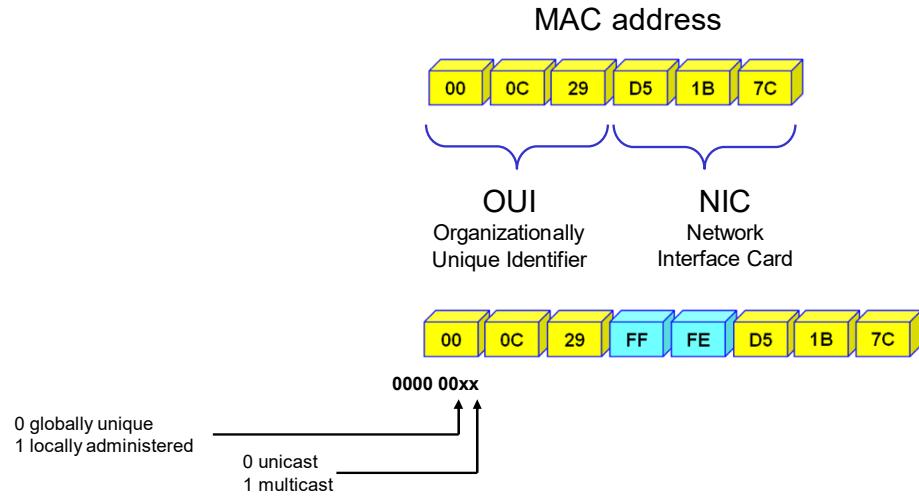
Si possono usare solo tra nodi dello stesso link

Indirizzi di sito: **FEC0 :: /48** (1111-1110-1100-0000)

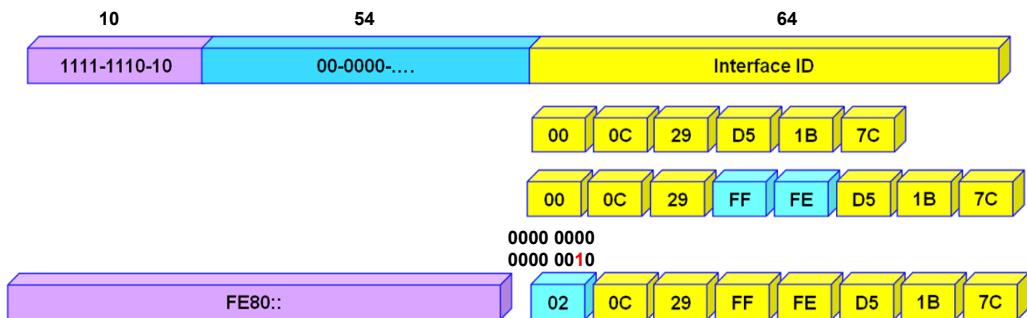


Si possono usare solo tra nodi di uno stesso sito. Sono sconsigliati da RFC 3879

Entrambi non possono essere usati su Internet.



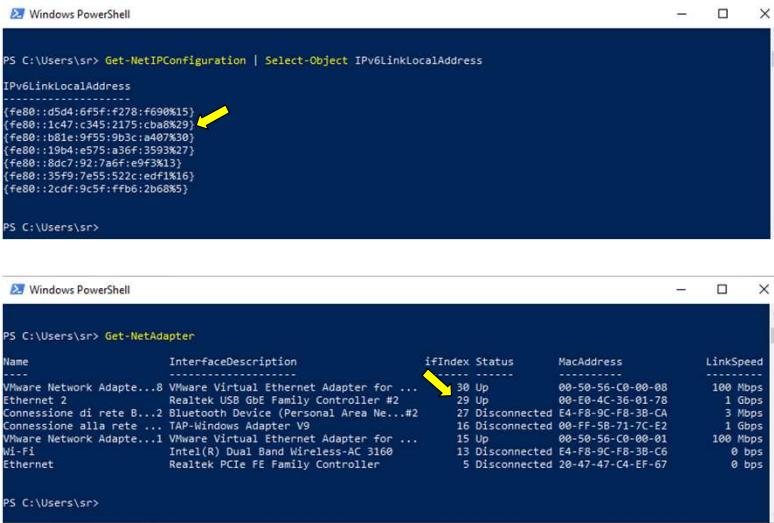
Indirizzi IPv6 EUI-64 (RFC 2373)



```
[root@LabReti ~]# ifconfig
eth0      Link encap:Ethernet  Hwaddr 00:0C:29:D5:1B:7C
          inet addr: 192.168.1.38  Bcast:192.168.1.255  Mask:255.255.255.0
          inet6 addr: fe80::20c:29ff:fed5:1b7c%64 Scope:Link
                     UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
                     RX packets:11 errors:0 dropped:0 overruns:0 frame:0
                     TX packets:13 errors:0 dropped:0 overruns:0 carrier:0
                     collisions:0 txqueuelen:1000
                     RX bytes:1421 (1.3 KiB)  TX bytes:1382 (1.3 KiB)
                     Interrupt:19 Base address:0x2000
```

IPv6 local link address non EUI-64

Windows OS non usa EUI-64 L'indirizzo IPv6 è calcolato partendo da un numero casuale. Manca il gruppo **FFFF** nella parte centrale.



```
PS C:\Users\sr> Get-NetIPConfiguration | Select-Object IPv6LinkLocalAddress
IPv6LinkLocalAddress
{fe80::d5d4:6f5f::278:f690%15}
{fe80::1c47:c345::2175:cb8x29} -----^
{fe80::b81e:9e55:9b3c:a407x30}
{fe80::19d4:e575:a36f:3593x27}
{fe80::8dc7:927a6f:ee9f3x13}
{fe80::35f9:7e55:522c:edf1x16}
{fe80::2cdf:9c5f:ffb6:2b68x5}

PS C:\Users\sr>

PS C:\Users\sr> Get-NetAdapter
Name           InterfaceDescription          ifIndex Status      MacAddress      LinkSpeed
----           InterfaceDescription          ifIndex Status      MacAddress      LinkSpeed
VMware Network Adapter...8 VMware Virtual Ethernet Adapter for ...    30 Up       00-50-56-C0-00-08   100 Mbps
Ethernet 2      Realtek USB GbE Family Controller #2        29 Up       00-E0-4C-36-01-78   1 Gbps
Connessione di rete ...2 Bluetooth Device (Personal Area Ne...#2    27 Disconnected E4-F8-9C-F8-3B-CA   3 Mbps
Connessione alla rete ... TAP-Windows Adapter V9                 16 Disconnected 00-FF-58-71-7C-E2   1 Gbps
VMware Network Adapter..1 VMware Virtual Ethernet Adapter for ...    15 Up       00-50-56-C0-00-01   100 Mbps
Wi-Fi          Intel(R) Dual Band Wireless-AC 3160            13 Disconnected E4-F8-9C-F8-3B-C6   0 bps
Ethernet        Realtek PCIe FE Family Controller             5 Disconnected 20-47-47-C4-EF-67   0 bps
```

Indirizzi Unicast (global)

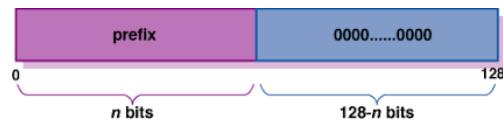
2xxx: **3xxx** :

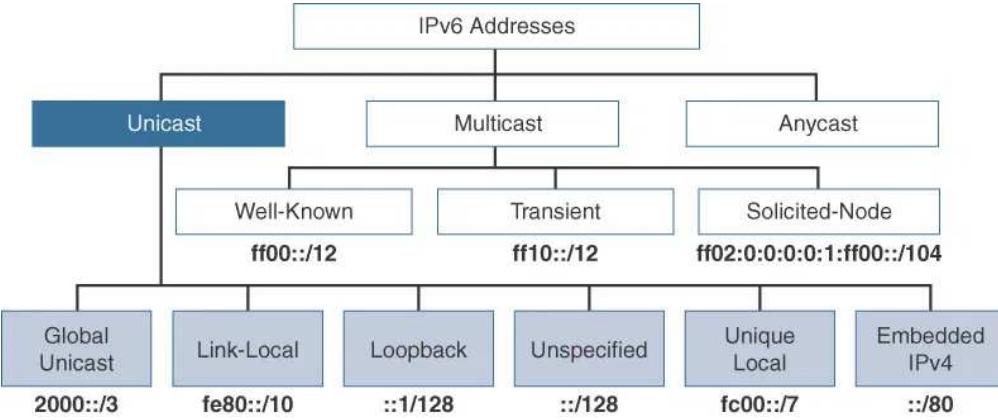
Indirizzi Multicast

FFxx:

Multicast Address	Scope	Group Within the Scope
FF01:0:0:0:0:0:1	Node-local	All-nodes address
FF01:0:0:0:0:0:2	Node-local	All-routers address
FF02:0:0:0:0:0:1	Link-local	All-nodes address
FF02:0:0:0:0:0:2	Link-local	All-routers address
FF02:0:0:0:0:0:5	Link-local	OSPF IGP
FF02:0:0:0:0:0:6	Link-local	OSPF IGP designated routers
FF02:0:0:0:0:0:D	Link-local	All PIM routers
FF02:0:0:0:0:0:16	Link-local	All MLDv2-capable routers
FF02:0:0:0:0:1:2	Link-local	All DHCP agents
FF02:0:0:0:1:FFXX:XX XX	Link-local	Solicited-node address
FF05:0:0:0:0:0:2	Site-local	All-routers address
FF05:0:0:0:0:1:3	Site-local	All DHCP servers

Indirizzi Anycast





IPv6 Ethernet encapsulation (RFC 2464)

I pacchetti IPv6 sono incapsulati in frame Ethernet esattamente come per i pacchetti IPv4: cambia solo il campo Ethertype (86DD invece che 0800).

NDP fornisce i seguenti servizi:

- Router discovery
- Prefix discovery
- Parameter discovery
- Address Autoconfiguration
- Address resolution
- Next-hop determination
- Neighbor unreachability detection
- Duplicate address detection
- Redirect messages

- Address Resolution Protocol (ARP),
- Internet Control Message Protocol (ICMP) router discovery,
- ICMP Redirect message

ARP non fa parte di ICMPv4. Aver incorporato il meccanismo di risoluzione degli indirizzi mac in ICMPv6, rende il tutto più indipendente dal tipo di mezzo fisico e consente l'utilizzo di meccanismi di sicurezza

NDP prevede 5 messaggi:

- Router Solicitation (ICMPv6 type 133)
- Router Advertisement (ICMPv6 type 134)
- Neighbor Solicitation (ICMPv6 type 135)
- Neighbor Advertisement (ICMPv6 type 136)
- Redirect (ICMPv6 type 137)

000 -111 = time insensitive (could be discarded)

1000 -1111 = priority (should not be discarded)

0 = uncharacterized

1 = filler (NetNews)

2 = unattended transfer (mail)

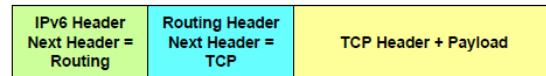
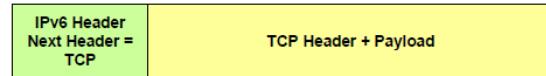
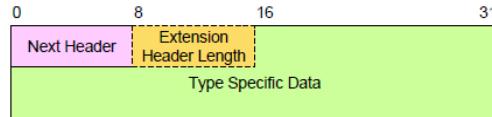
4 = bulk (ftp)

6 = interactive (telnet)

7 = Internet control

8 = video

15 = low quality audio

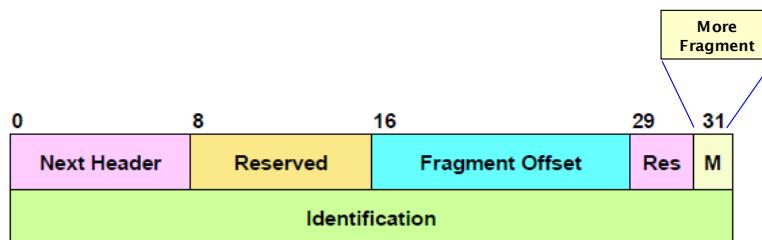


0	HBH	Hop-by-Hop option (IPv6) (Jumbogram)
1	ICMP	Internet Control Message (IPv4)
2	IGMP	Internet Group Management (IPv4)
3	GGP	Gateway-to-Gateway Protocol
4	IP	IP in IP (IPv4 encapsulation)
5	ST	Stream
6	TCP	Transmission Control
8	EGP	Exterior Gateway Protocol
9	IGP	Any private interior gateway
16	CHAOS	Chaos
17	UDP	User Datagram
29	ISO-TP4	ISO Transport Protocol Class 4
36	XTP	XTP
43	RH	Routing header (IPv6)
44	FH	Fragmentation header (IPv6)
45	IDRP	Inter-Domain Routing Protocol
46	RSVP	Reservation Protocol
50	ESP	Encapsulating Security Payload
51	AH	Authentication header (IPv6)
54	NHRP	NBMA Next Hop Resolution Protocol
58	ICMP	Internet Control Message (IPv6)
59	Null	No next header (IPv6)
60	DOH	Destination Options header (IPv6)
80	ISO-IP	ISO Internet Protocol (CLNP)
83	VINES	VINES
88	IGRP	IGRP
89	OSPF	OSPF (Open Shortest Path First)
93	AX.25	AX.25 Frames

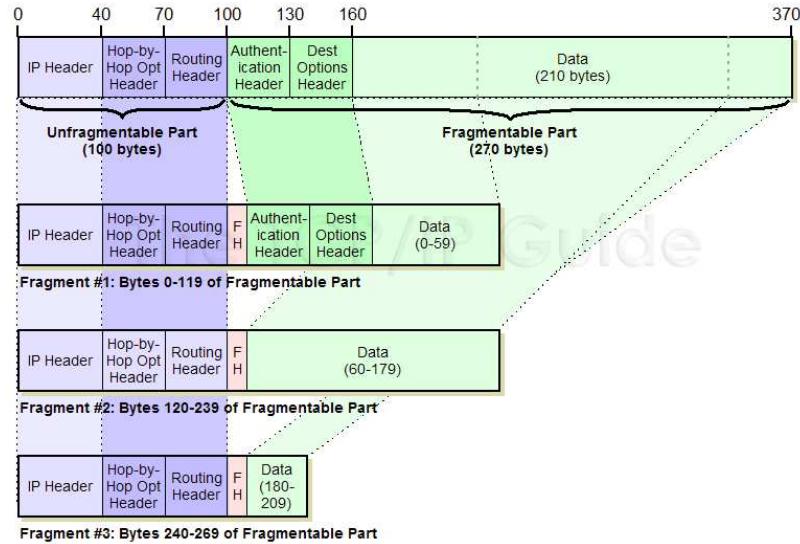
In IPv4 qualunque nodo può frammentare.

In IPv6 la frammentazione è scoraggiata. (RFC 1981 – MTU Path discovery for IPv6). Il minimo MTU è 1280 byte, con 1500 raccomandati.

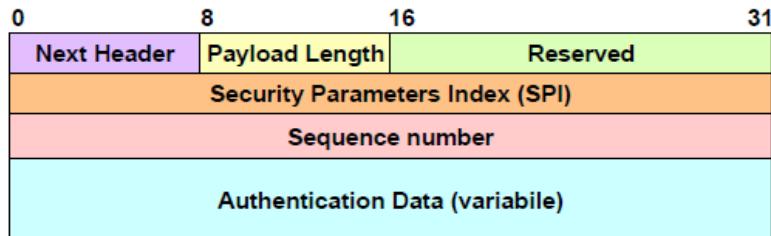
Solo il mittente può frammentare, tramite un Fragment Header.



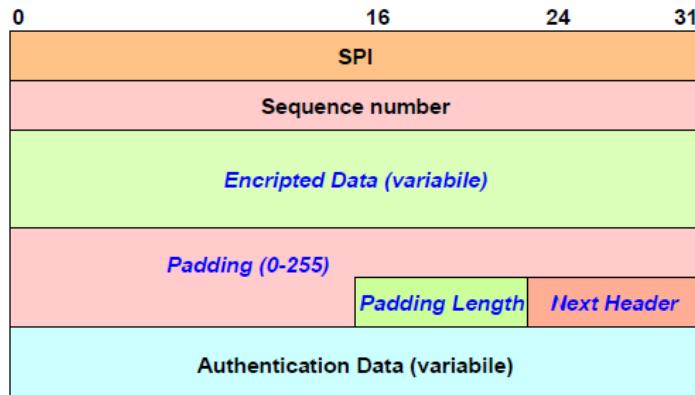
Il pacchetto IPv6 può essere frammentato solo dopo il FH. Tutto ciò che è prima non deve essere frammentato.



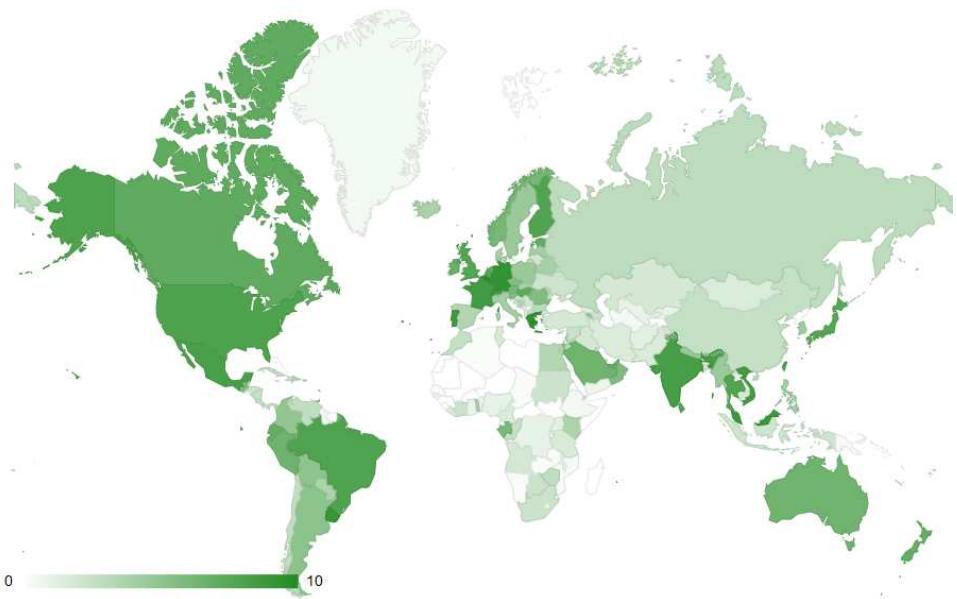
Authentication Header Garantisce l'autenticità e la correttezza del pacchetto.



Encrypted security payload Header. Fornisce la cifratura del payload.



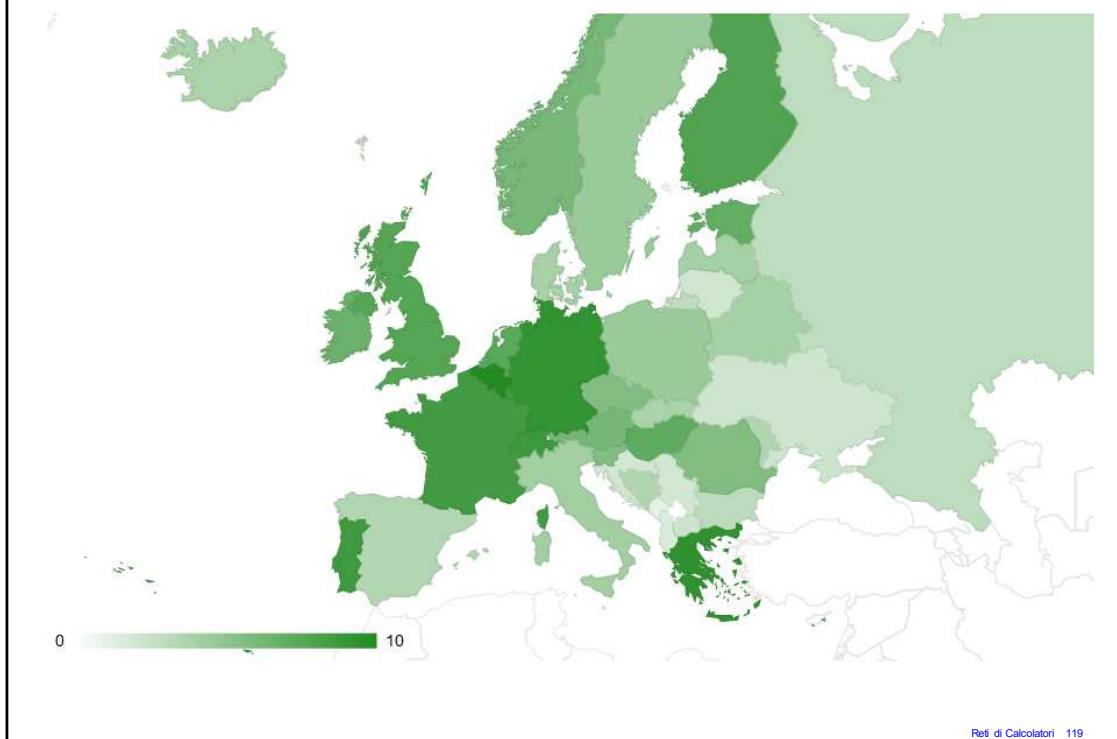
Diffusione di IPv6



<https://6lab.cisco.com/stats/index.php>

Reti di Calcolatori 118

Diffusione di IPv6 in Europa

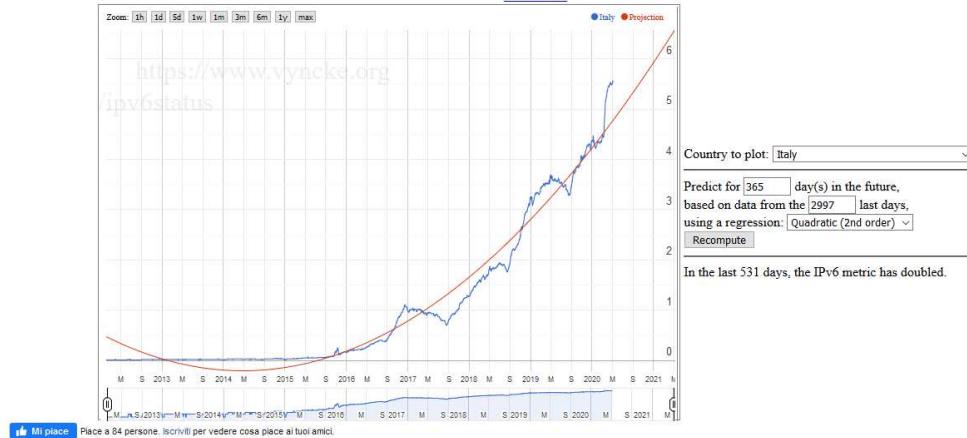


Reti di Calcolatori 119

Projection of IPv6 %-age of IPv6-Enabled Web Browsers (courtesy Google) in Italy

%-age of IPv6-Enabled Web Browsers (courtesy Google)

Metric to display: [allocated prefixes](#) - [announced prefixes](#) - [alive prefixes](#) - [global routing table](#) - [IPv6 web browsers \(Google\)](#) - [IPv6 web browsers \(APnic\)](#) - [IPv6 web browsers \(Akamai\)](#) - [IPv6 web servers](#)

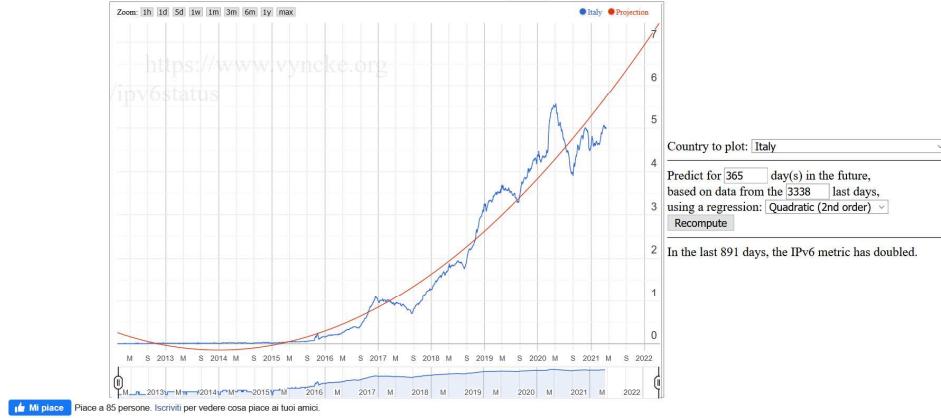


Thanks to Google for their Chart tools, to Erik Kline for his suggestion and to [Tom-Alexander](#) for his polynomial regression package.

Projection of IPv6 %-age of IPv6-Enabled Web Browsers (courtesy Google) in Italy

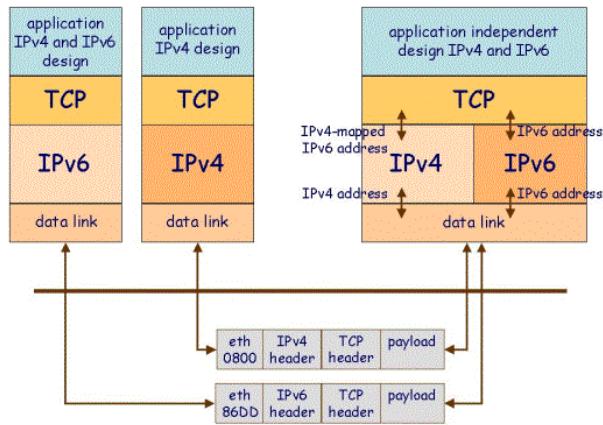
%-age of IPv6-Enabled Web Browsers (courtesy Google)

Metric to display: [allocated prefixes](#) - [announced prefixes](#) - [alive prefixes](#) - [global routing table](#) - [IPv6 web browsers \(Google\)](#) - [IPv6 web browsers \(APnic\)](#) - [IPv6 web browsers \(Akamai\)](#) - [IPv6 web servers](#)



Thanks to Google for their Chart tools, to Erik Kline for his suggestion and to [Tom-Alexander](#) for his polynomial regression package.

L'approccio dual-stack

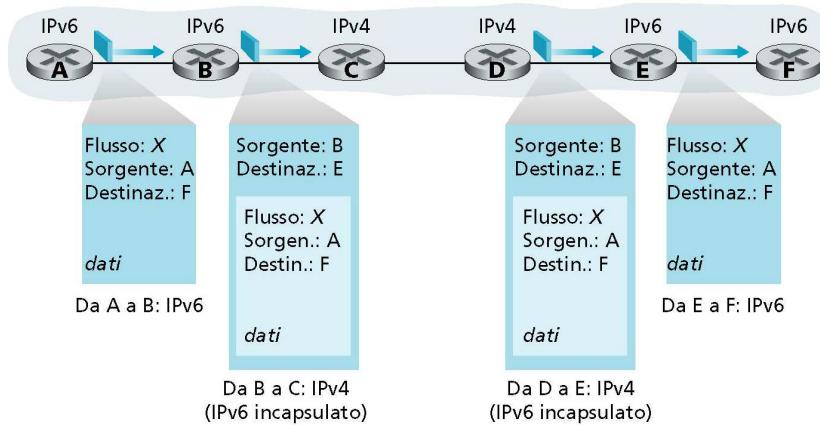


Tunneling

Vista logica

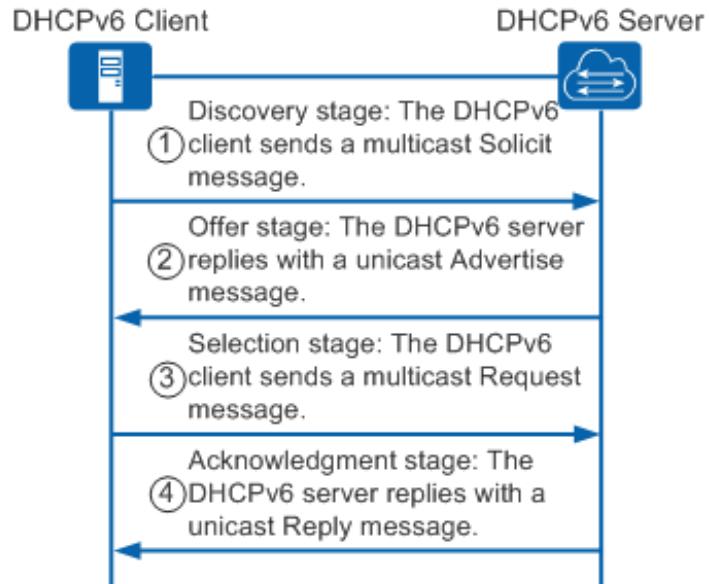


Vista fisica

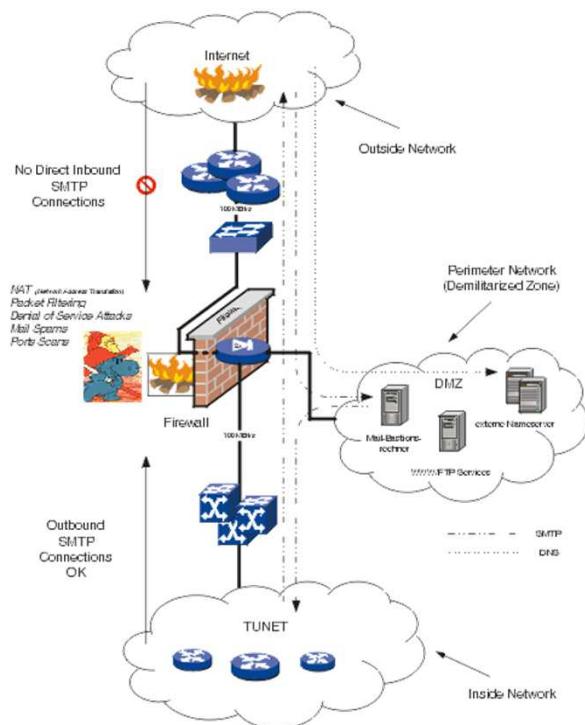


ICMP Tipo	Codice	Deserzione
0	0	risposta al messaggio di eco (a ping) - <i>echo replay</i>
3	0	rete di destinazione irraggiungibile - <i>destination network unreachable</i>
3	1	host di destinazione irraggiungibile - <i>destination host unreachable</i>
3	2	protocollo di destinazione irraggiungibile - <i>destination protocol unreachable</i>
3	3	porta di destinazione irraggiungibile - <i>destination port unreachable</i>
3	6	rete di destinazione sconosciuta - <i>destination network unknown</i>
3	7	host di destinazione sconosciuto - <i>destination host unreachable</i>
4	0	strozzamento della sorgente (controllo della congestione) - <i>source quench</i>
8	0	richiesta di eco - <i>echo request</i>
9	0	annuncio dal router - <i>router advertisement</i>
10	0	scoperta del router - <i>router discovery</i>
11	0	TTL scaduto - <i>TTL expired</i>
12	0	cattiva intestazione IP - <i>IP header bad</i>

Type	Meaning
1	Destination Unreachable
2	Packet Too Big
3	Time Exceeded
4	Parameter Problem
128	Echo Request
129	Echo Reply
130	Group Membership Query
131	Group Membership Report
132	Group Membership Reduction
133	Router Solicitation
134	Router Advertisement
135	Neighbor Solicitation
136	Neighbor Advertisement
137	Redirect
138	Router Renumbering

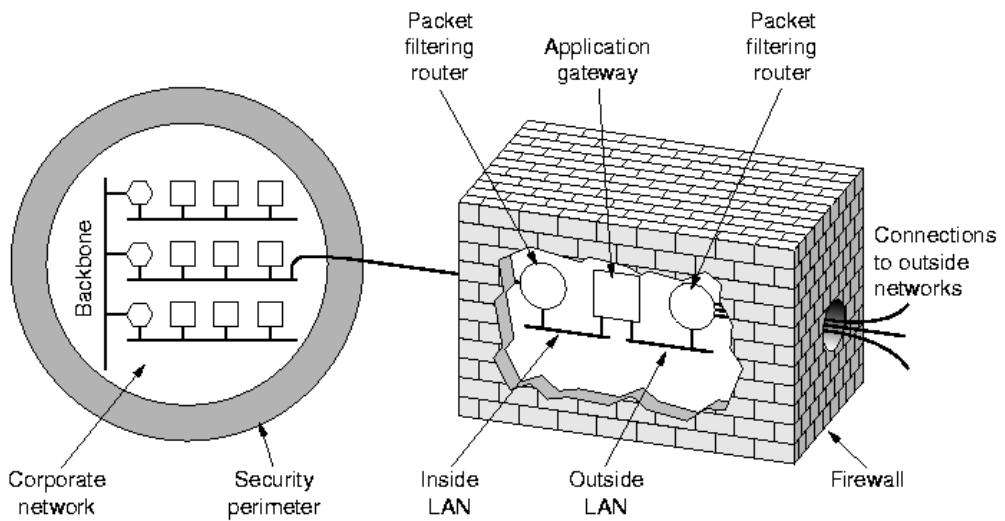


Firewall

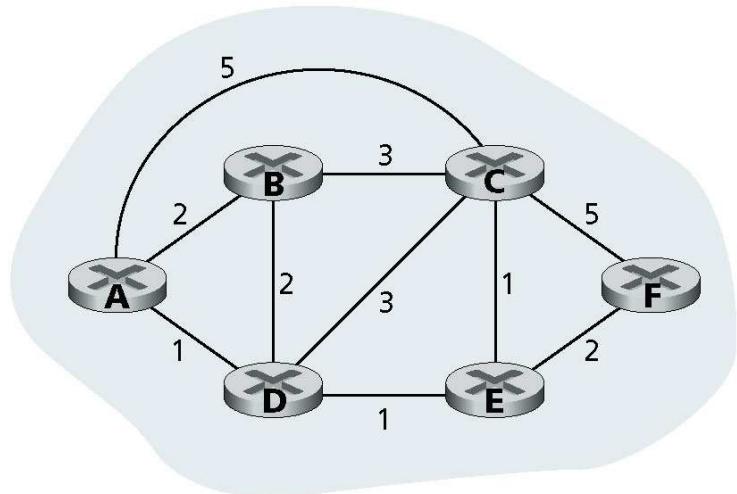


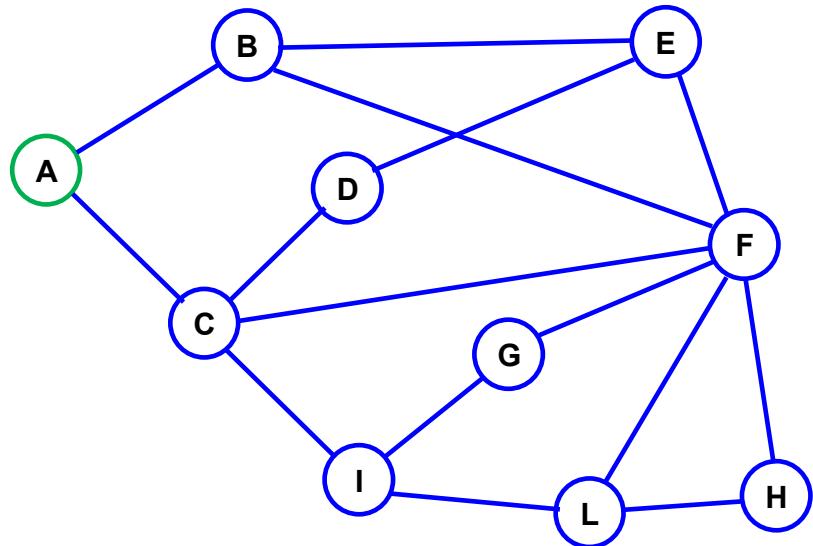
Reti di Calcolatori 127

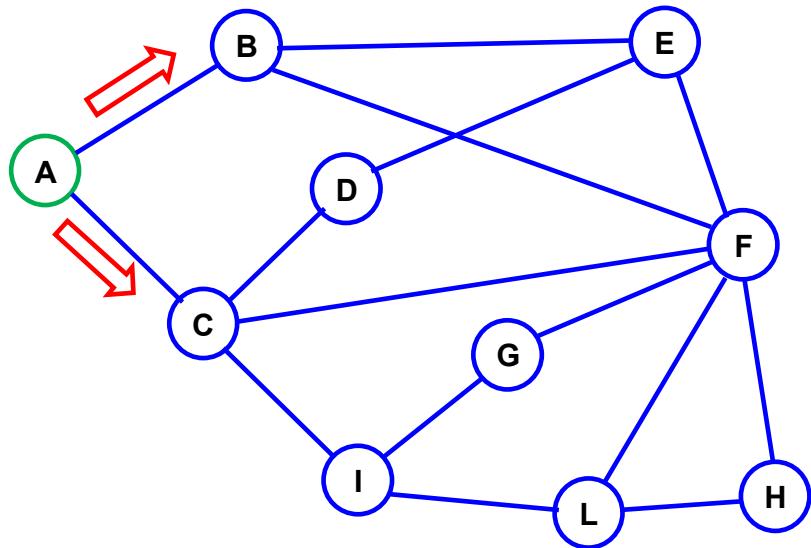
Un firewall è un sistema di sicurezza della rete informatica che limita il traffico Internet in entrata, in uscita o all'interno di una rete privata. Questo software o unità hardware-software dedicata funziona bloccando o consentendo in maniera selettiva i pacchetti di dati.

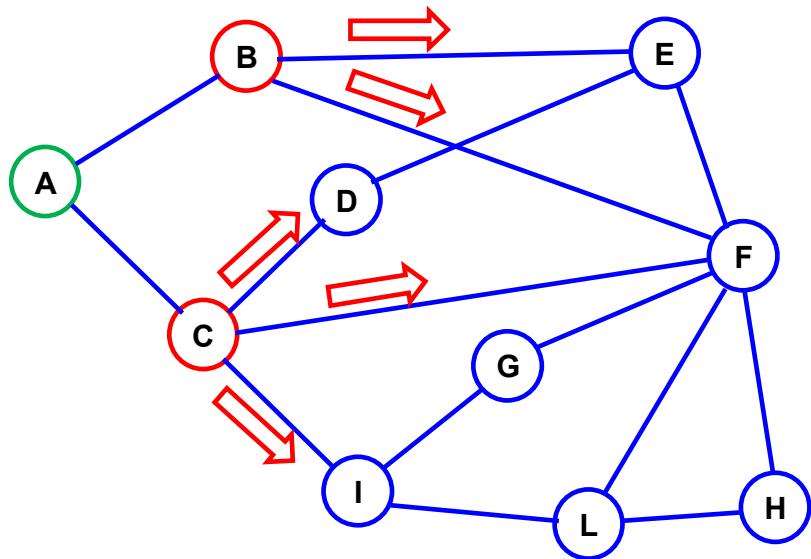


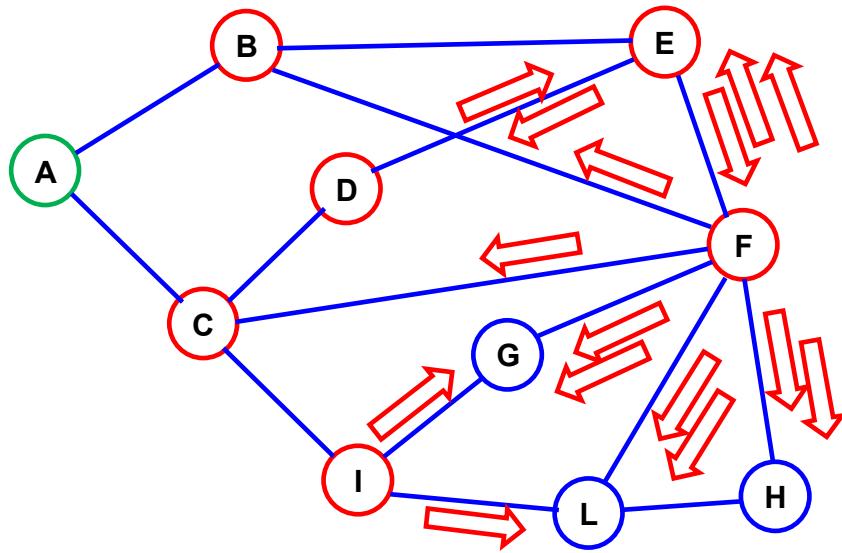
- Statici
 - Dinamici
-
- Locali
 - Globali

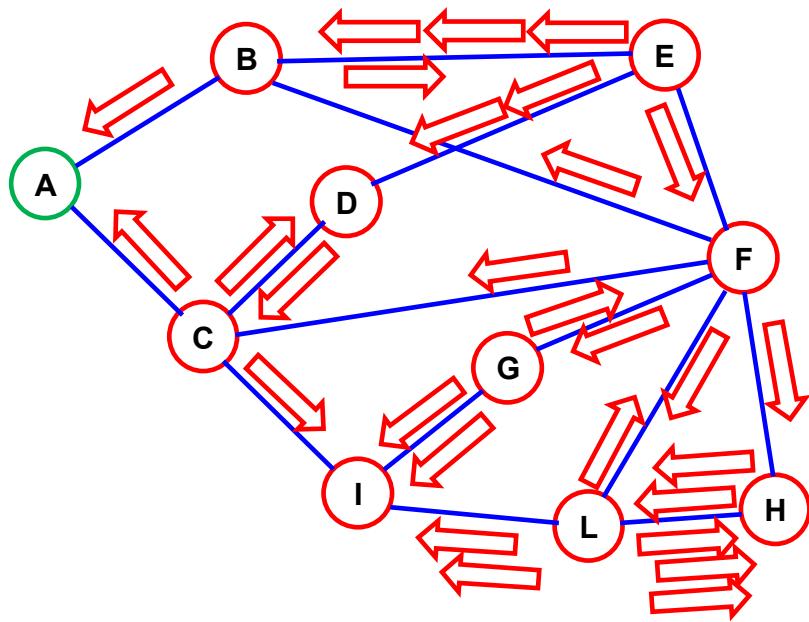


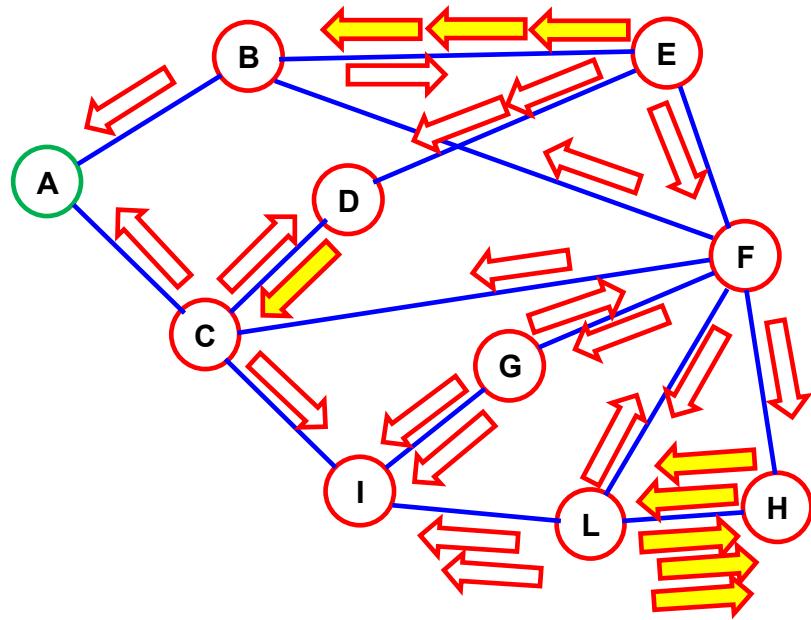












Vantaggi:

- Trova un percorso verso la destinazione, se esiste.
- Trova il percorso migliore.

Svantaggi:

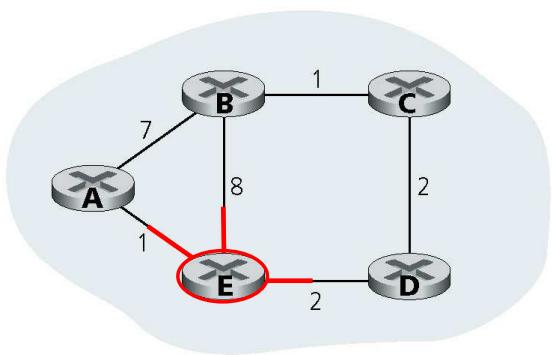
- Crea un numero elevato di pacchetti duplicati, saturando la rete.

Il flooding deve essere controllato in qualche modo.

Controllo del flooding

- 1) Usare un ID per i pacchetti
- 2) Usare un HOP limit
- 3) Eliminare i cicli dal grafo
- 4) Inoltrare solo lungo uno Spanning Tree

Distance Vectors

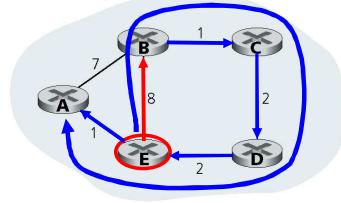


costo verso la destinazione via

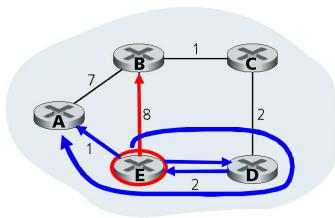
$D^E()$	A	B	D
A	1	14	5
B	7	8	5
C	6	9	4
D	4	11	2

destinazione

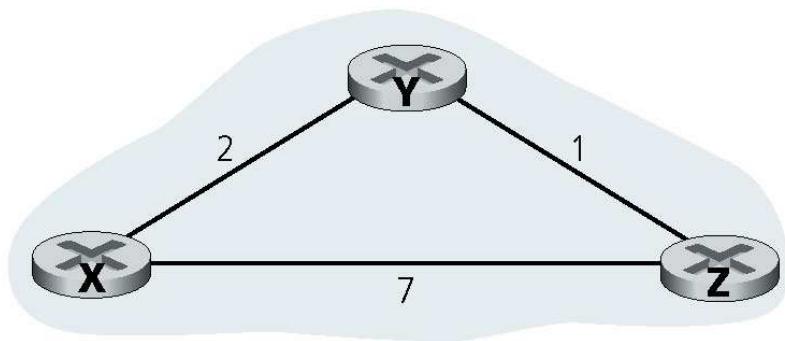
Distance Vectors



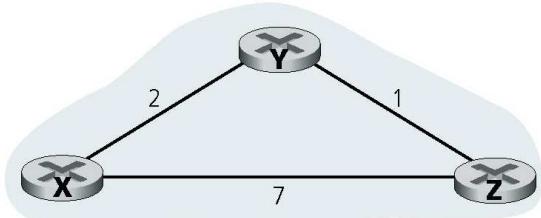
	A	B	D
A	1	14	5
B	7	8	5
C	6	9	4
D	4	11	2



	A	B	D
A	1	14	5
B	7	8	5
C	6	9	4
D	4	11	2



Distance Vectors



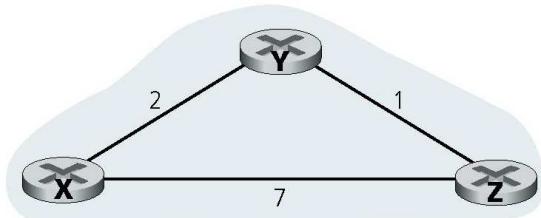
Uscita			
Destinazione	X	Y	Z
X	∞	∞	∞
Y	∞	∞	∞
Z	∞	∞	∞

X	Y	Z
Y	∞	∞
Z	∞	∞

Y	X	Z
X	∞	∞
Z	∞	∞

Z	X	Y
X	∞	∞
Y	∞	∞

Distance Vectors



Uscita			
Destinazione	X	Y	Z
X	∞	∞	∞
Y	∞	∞	∞
Z	∞	∞	∞

X	Y	Z
Y	∞	∞
Z	∞	∞
X	∞	∞
Y	∞	∞
Z	∞	∞

X	Y	Z
Y	2	∞
Z	∞	7
X	∞	∞
Y	∞	∞
Z	∞	∞

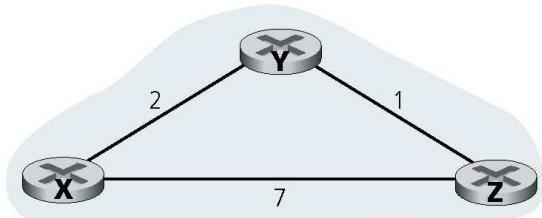
X
Y
Z
X
Y
Z

Distance Vectors

Destinazione

Uscita

X	Y	Z
Y	∞	∞
Z	∞	∞



X	Y	Z
Y	∞	∞
Z	∞	∞

X	Y	Z
Y	2	∞
Z	∞	7

X	
Y	2
Z	7

X	Y	Z
Y	2	8
Z	3	7

Y	X	Z
X	∞	∞
Z	∞	∞

Y	X	Z
X	2	∞
Z	∞	1

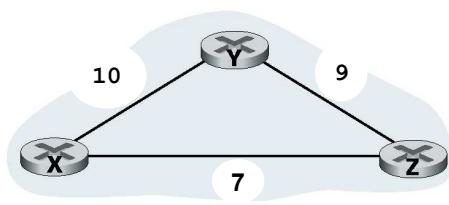
Z	X	Y
X	∞	∞
Y	∞	1

Y	
X	2
Z	1

Z	
X	7
Y	1

Y	X	Z
X	2	8
Z	9	1

Z	X	Y
X	7	3
Y	9	1



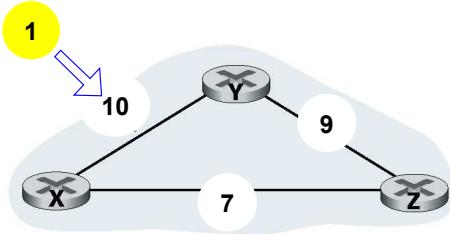
Uscita		
X	Y	Z
Y	∞	∞
Z	∞	∞

X	Y	Z	X	Y	Z
Y	10	∞	Y	10	16
Z	∞	7	Z	19	7
Y	X	Z	Y	X	Z
X	10	∞	X	10	16
Z	∞	9	Z	17	9
Z	X	Y	Z	X	Y
X	7	∞	X	7	19
Y	∞	9	Y	17	9

Distance Vectors: miglioramenti

Uscita

X	Y	Z
Y	∞	∞
Z	∞	∞



X	Y	Z
Y	10	16
Z	19	7

X	Y	Z
Y	1	16
Z	19	7

Y	X	Z
X	10	16
Z	17	9

Y	X	Z
X	1	16
Z	17	9

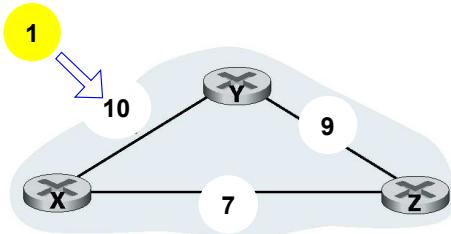
Z	X	Y
X	7	19
Y	17	9

Z	X	Y
X	7	19
Y	17	9

Per Z:
uscita X verso Y

Distance Vectors: miglioramenti

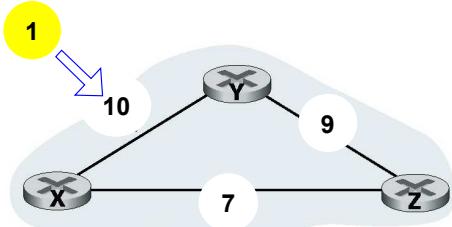
Uscita		
X	Y	Z
Y	∞	∞
Z	∞	∞



X	Y	Z	X	Y	Z	X	Y	Z	X	Y	Z
Y	10	16	Y	1	16	Y	1		Y	1	16
Z	19	7	Z	19	7	Z	7		Z	10	7
Y	X	Z	X	X	Z	X	Y		Y	X	Z
X	10	16	X	1	16	X	1		X	1	16
Z	17	9	Z	17	9	Z	9		Z	8	9
Z	X	Y	Z	X	Y	Z	X	7	Z	X	Y
X	7	19	X	7	19	X	7		X	7	10
Y	17	9	Y	17	9	Y	9		Y	8	9

Distance Vectors: miglioramenti

Uscita		
X	Y	Z
Y	∞	∞
Z	∞	∞



X	Y	Z
Y	10	16
Z	19	7

X	Y	Z
Y	1	16
Z	19	7

X	
Y	1
Z	7

X	Y	Z
Y	1	16
Z	10	7

Y	X	Z
X	10	16
Z	17	9

Y	X	Z
X	1	16
Z	17	9

Y	
X	1
Z	9

Y	X	Z
X	1	16
Z	8	9

Z	X	Y
X	7	19
Y	17	9

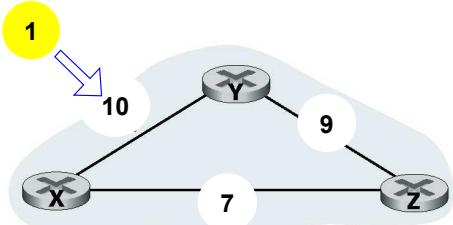
Z	X	Y
X	7	19
Y	17	9

Z	
X	7
Y	9

Z	X	Y
X	7	10
Y	8	9

Distance Vectors: miglioramenti

Uscita		
X	Y	Z
Y	∞	∞
Z	∞	∞



X	Y	Z
Y	10	16
Z	19	7

X	Y	Z
Y	1	16
Z	19	7

X	
Y	1
Z	7

X	Y	Z
Y	1	16
Z	10	7

X	
Y	1
Z	7

Y	X	Z
X	10	16
Z	17	9

Y	X	Z
X	1	16
Z	17	9

Y	
X	1
Z	9

Y	X	Z
X	1	16
Z	8	9

Y	
X	1
Z	8

Z	X	Y
X	7	19
Y	17	9

Z	X	Y
X	7	19
Y	17	9

Z	
X	7
Y	9

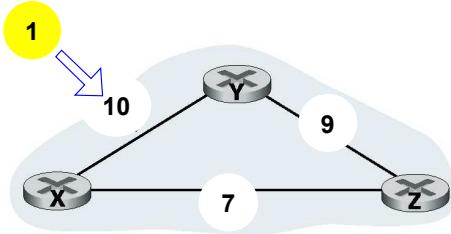
Z	X	Y
X	7	10
Y	8	9

Z	
X	7
Y	8

Reti di Calcolatori 148

Distance Vectors: miglioramenti

Uscita		
X	Y	Z
Y	∞	∞
Z	∞	∞



X	Y	Z
Y	1	16
Z	10	7

X	
Y	1
Z	7

X	Y	Z
Y	1	15
Z	9	7

Y	X	Z
X	1	16
Z	8	9

Y	
X	1
Z	8

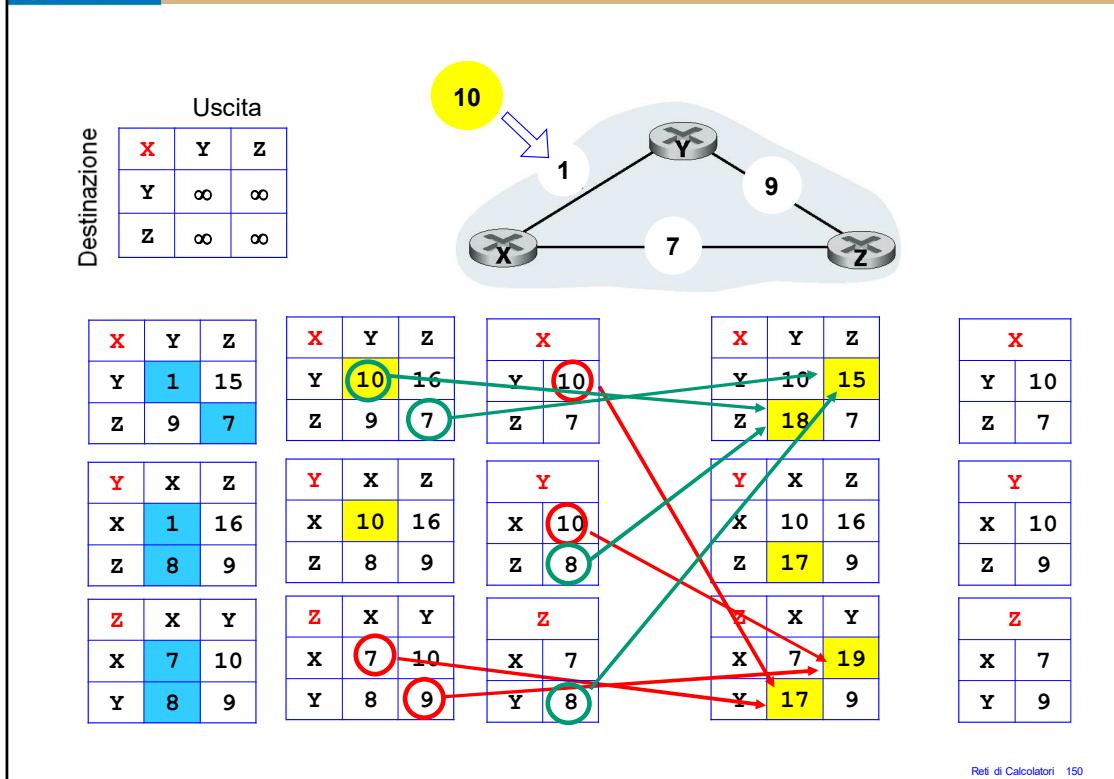
X	Y	Z
Y	1	16
Z	8	9

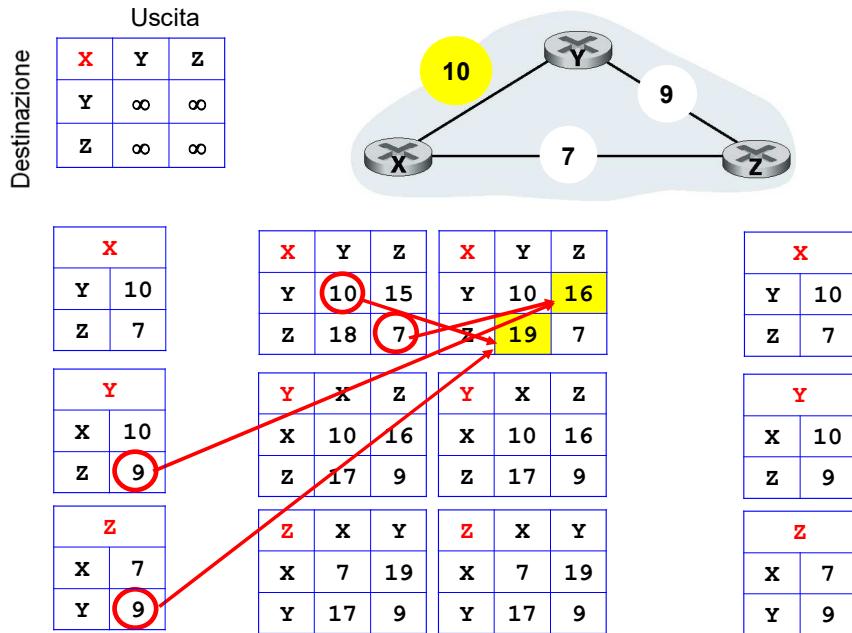
Z	X	Y
X	7	10
Y	8	9

Z	
X	7
Y	8

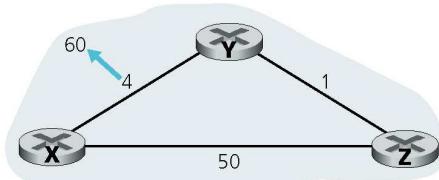
Z	X	Y
X	7	10
Y	8	9

Distance Vectors: peggioramenti





Distance Vectors



Miglior percorso noto

via		
D^Y	X	Z
X	4	6

via		
D^Y	X	Z
X	60	6

Miglior percorso noto

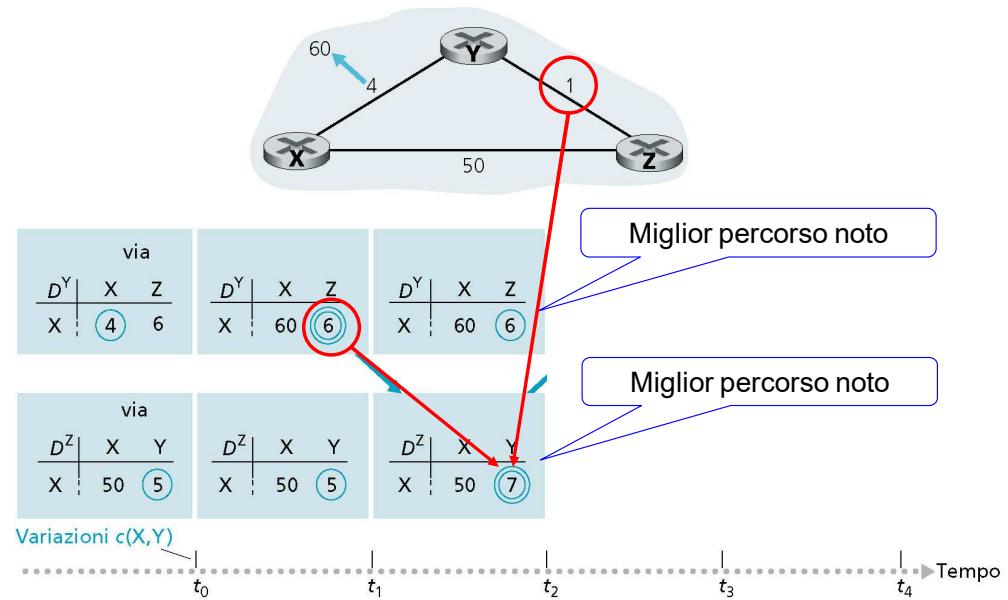
via		
D^Z	X	Y
X	50	5

via		
D^Z	X	Y
X	50	5

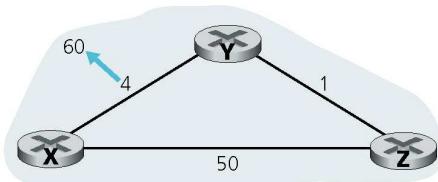
Variazioni $c(X,Y)$



Distance Vectors



Conteggio all'infinito



via		
D^Y	X	Z
X	4	6

via		
D^Y	X	Z
X	60	6

via		
D^Y	X	Z
X	60	6

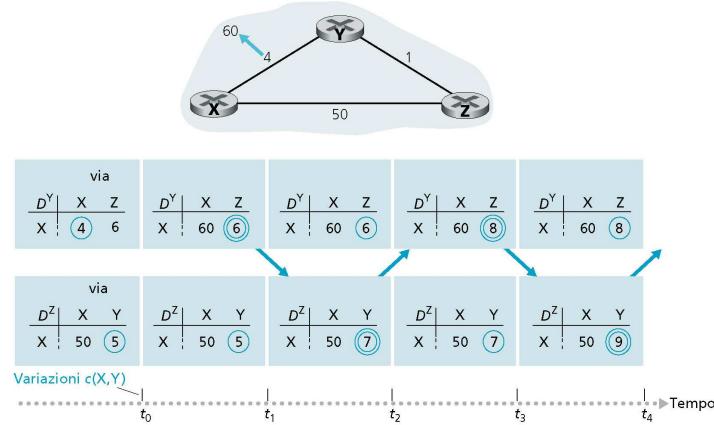
via		
D^Y	X	Z
X	60	8

via		
D^Y	X	Z
X	60	8

Variazioni $c(X,Y)$



Poisoned reverse



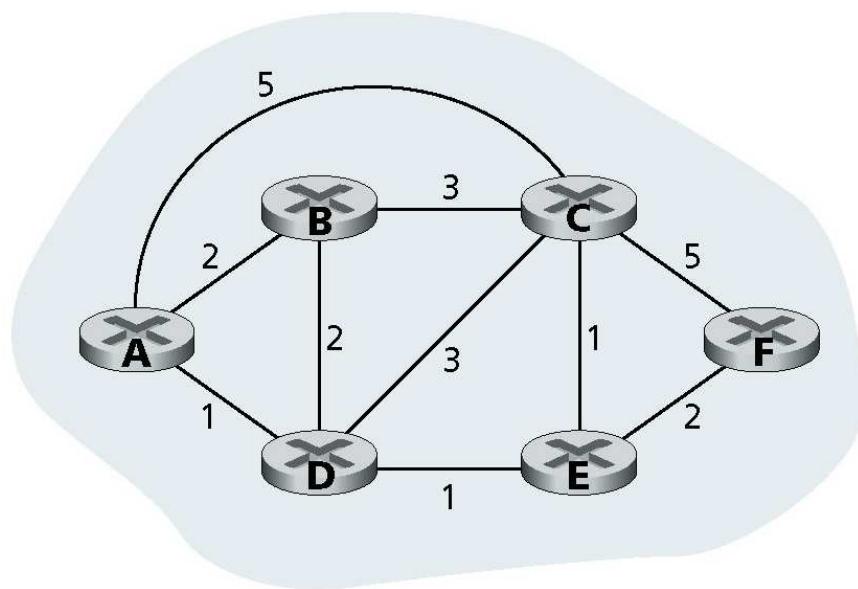
Possibile soluzione: inviare ∞ al nodo che si ha come uscita migliore per raggiungere un altro nodo

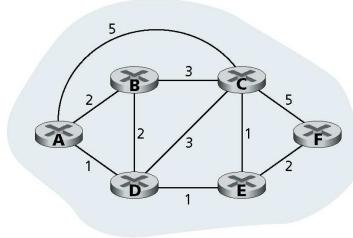
(es: z dice a y che la sua distanza verso x è ∞ se il suo percorso migliore verso x parte da y).

Reti di Calcolatori 155

Poison Reverse è un metodo di evitamento del loop utilizzato nel Routing Information Protocol (RIP) nel Distance Vector Routing Protocol (DVRP) che consente a un'interfaccia abilitata RIP di impostare il costo del percorso (cioè impostare il conteggio degli hop che generalmente varia da 0 a 15) che ottiene dai router vicini a 16 (qui 16 sarà considerato infinito) che mostra che un determinato percorso non è raggiungibile, quindi rimanda indietro il percorso. Quando il percorso adiacente riceve questo percorso con costo 16, elimina i percorsi inutili dalla sua tabella di routing che impedisce ulteriormente il looping.

Poison Reverse è usato per affrontare i problemi di conteggio all'infinito e si può immaginare come un inverso del metodo Split Horizon. Con l'aiuto del veleno inverso, possiamo pubblicizzare la pubblicità del percorso che sarebbe soppressa da un orizzonte diviso con una distanza di infinito.



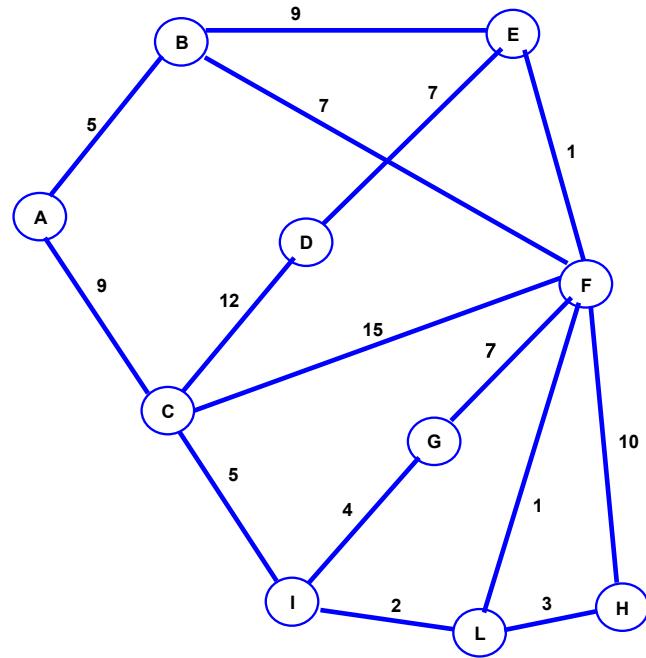


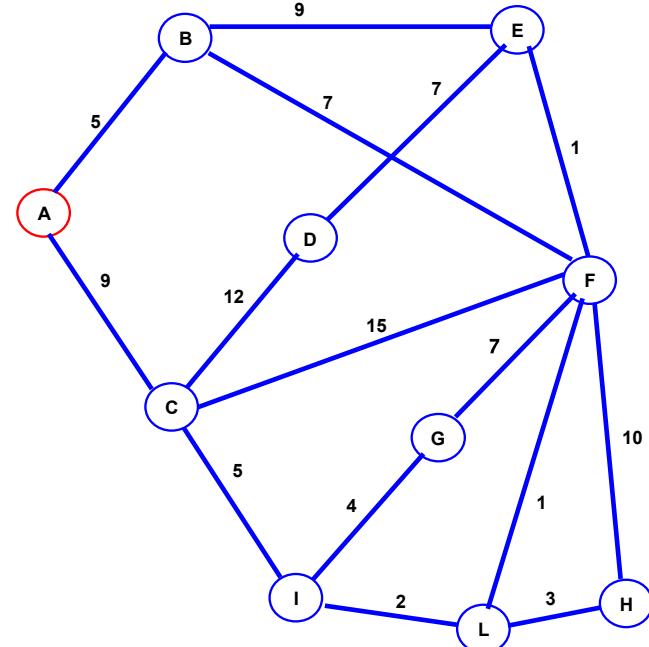
La prima versione di Arpanet utilizzava il **Distance Vector Routing**.

L'algoritmo venne sostituito dopo il 1979 con il **Link State Routing** in quanto il primo non teneva conto delle capacità dei canali ma solo dei tempi di ritardo di trasmissione.

Con questo algoritmo, ogni router indaga sui propri vicini e determina i **costi** di comunicazione. Spedisce quindi le informazioni in broadcast, in modo che ogni router possa calcolare i percorsi ottimali.

Il calcolo è basato sull'algoritmo di **Dijkstra**.

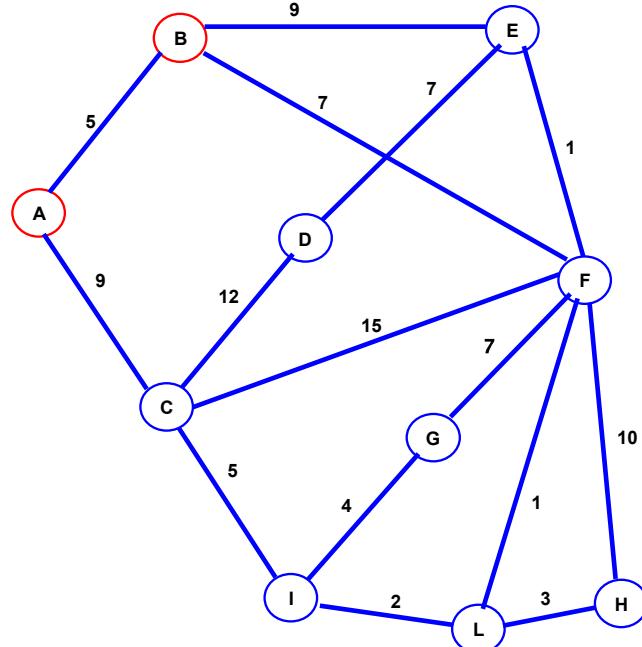




A

B = 5

C = 9



A

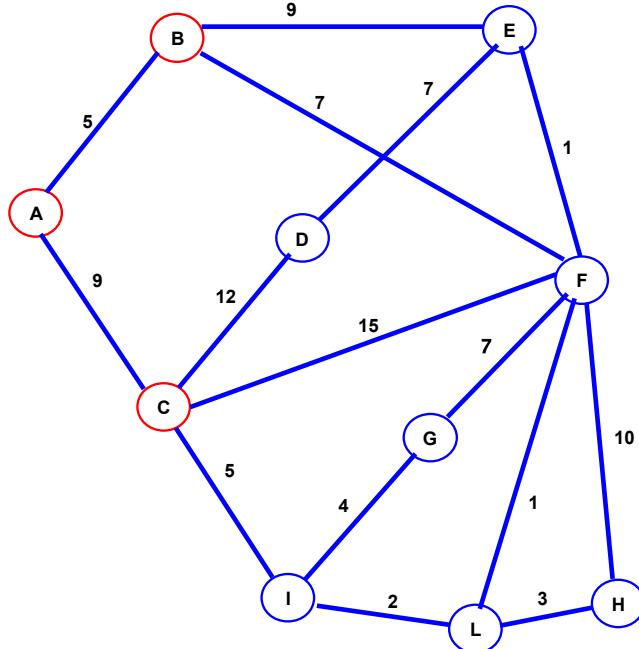
$$B = 5$$

$$C = 9$$

$$E = 5+9 = 14$$

$$F = 5+7= 12$$

Dijkstra



A

B = 5

C = 9

$$E = 5+9 = 14$$

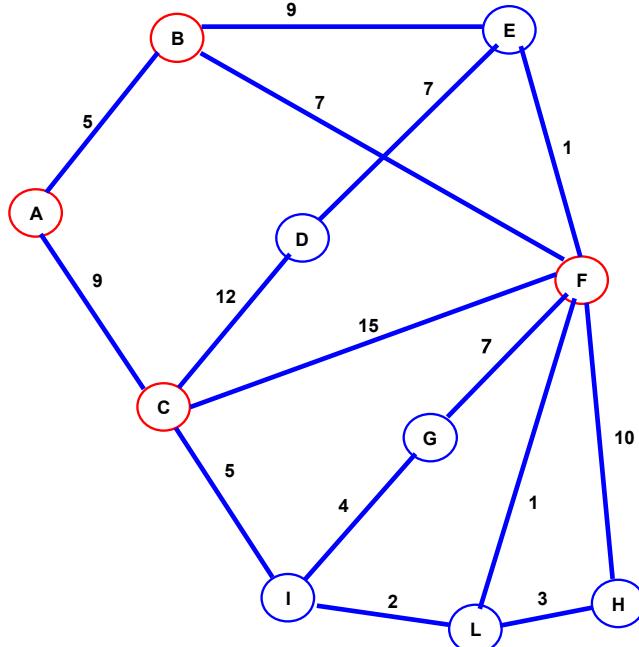
$$F = 5+7= 12$$

$$D = 9+12= 21$$

$$F = 9+15= 24$$

$$I = 9+5 = 14$$

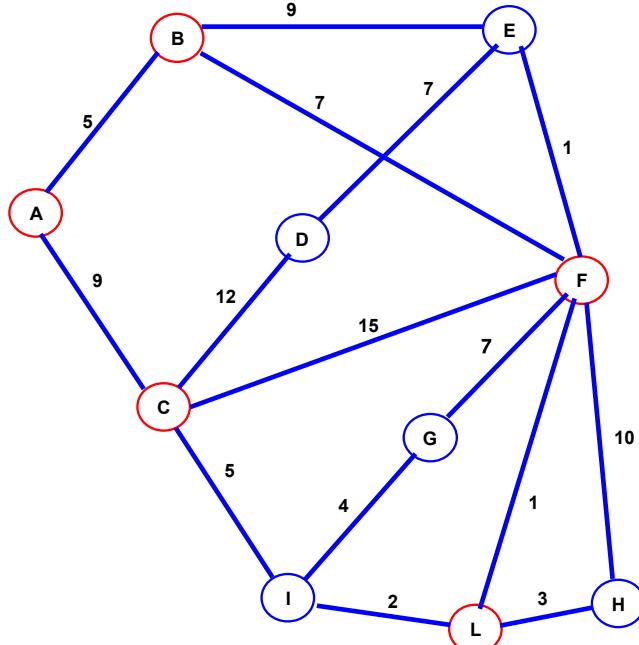
Dijkstra



A
 $B = 5$
 $C = 9$
 $F = 12$

$E = 5+9 = 14$
 $D = 9+12 = 21$
 $I = 9+5 = 14$
 $E = 12+1 = 13$
 $G = 12+7 = 19$
 $H = 12+10 = 22$
 $L = 12+1 = 13$

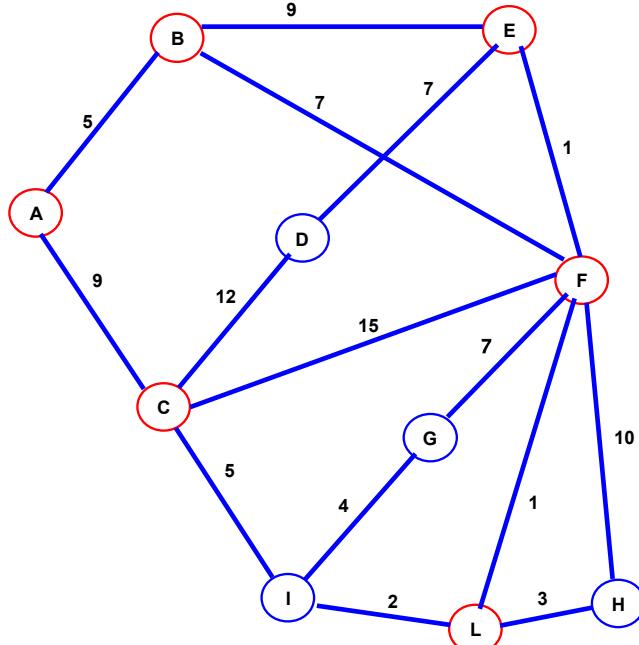
Dijkstra



A
 $B = 5$
 $C = 9$
 $F = 12$
 $L = 13$

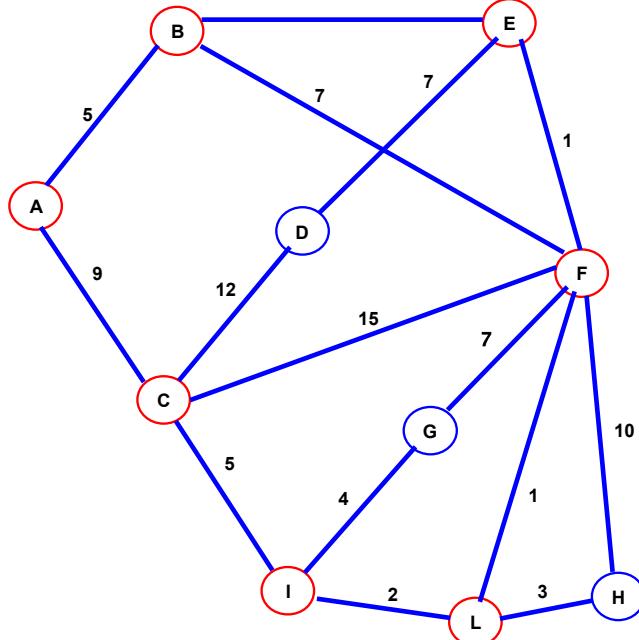
$E = 5+9 = 14$
 $D = 9+12 = 21$
 $I = 9+5 = 14$
 $E = 12+1 = 13$
 $G = 12+7 = 19$
 $H = 12+10 = 22$
 $I = 13+2 = 15$
 $H = 13+3 = 16$

Dijkstra



A
 B = 5
 C = 9
 F = 12
 L = 13
 E = 13

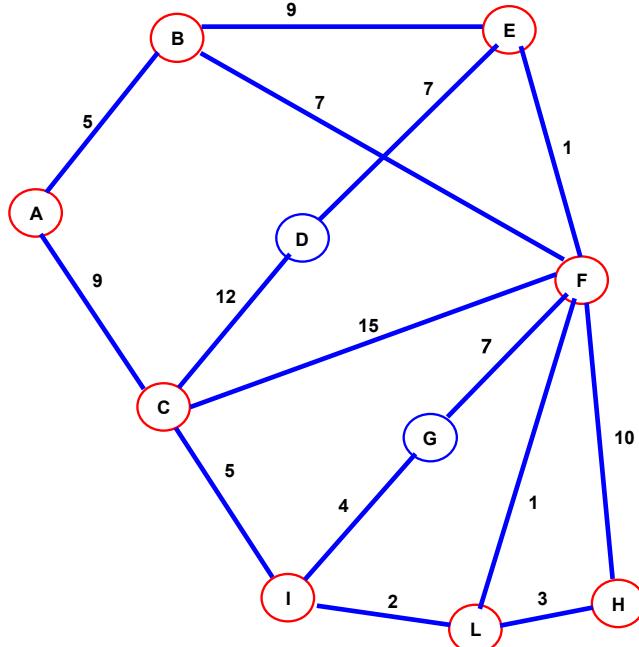
D = 9+12= 21
 I = 9+5 = 14
 G = 12+7 = 19
 H = 12+10 = 22
 I = 13+2 = 15
 H = 13+3 = 16
 D = 13+7 = 20



A
 $B = 5$
 $C = 9$
 $F = 12$
 $L = 13$
 $E = 13$
 $I = 14$

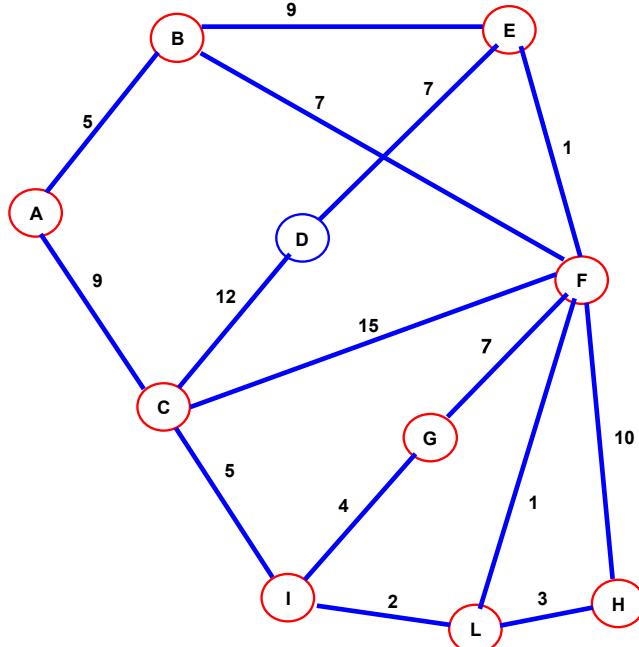
$D = 9+12= 21$
 $G = 12+7 = 19$
 $H = 12+10 = 22$
 $H = 13+3 = 16$
 $D = 13+7 = 20$
 $G = 14+4 = 18$

Dijkstra



A
 $B = 5$
 $C = 9$
 $F = 12$
 $L = 13$
 $E = 13$
 $I = 14$
 $H = 16$

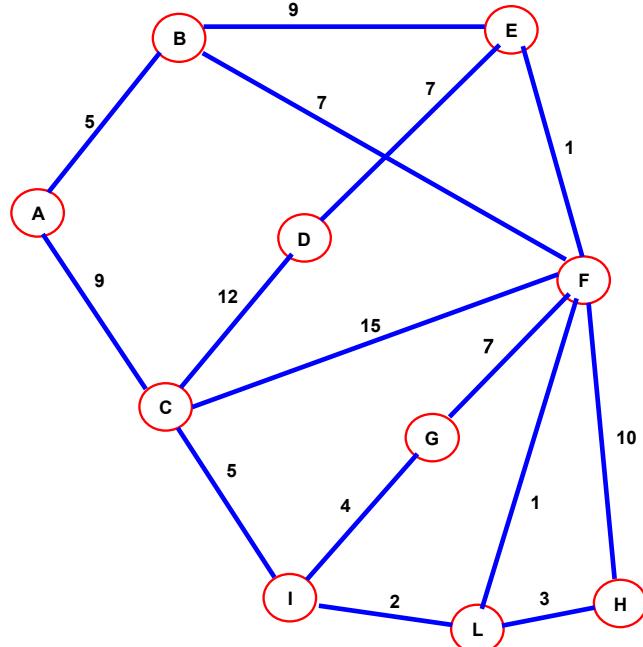
$$\begin{aligned}
 D &= 9+12= 21 \\
 G &= 12+7 = 19 \\
 D &= 13+7 = 20 \\
 G &= 14+4 = 18
 \end{aligned}$$



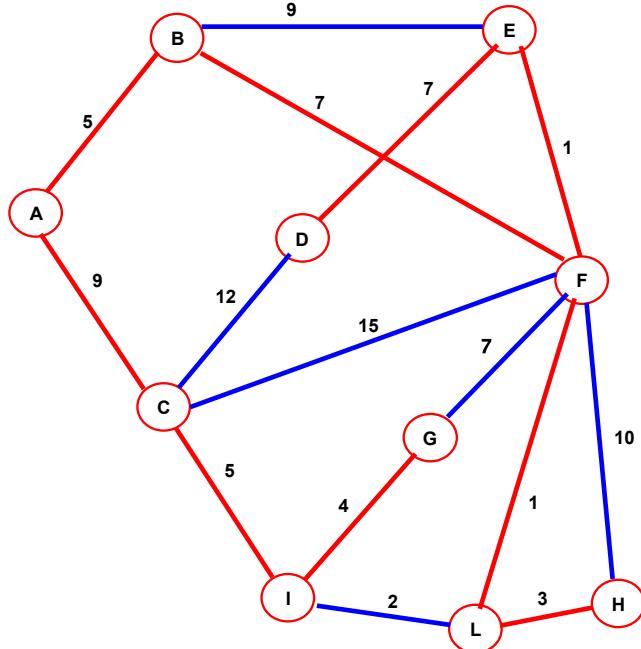
A
 B = 5
 C = 9
 F = 12
 L = 13
 E = 13
 I = 14
 H = 16
 G = 18

$$D = 9+12= 21$$

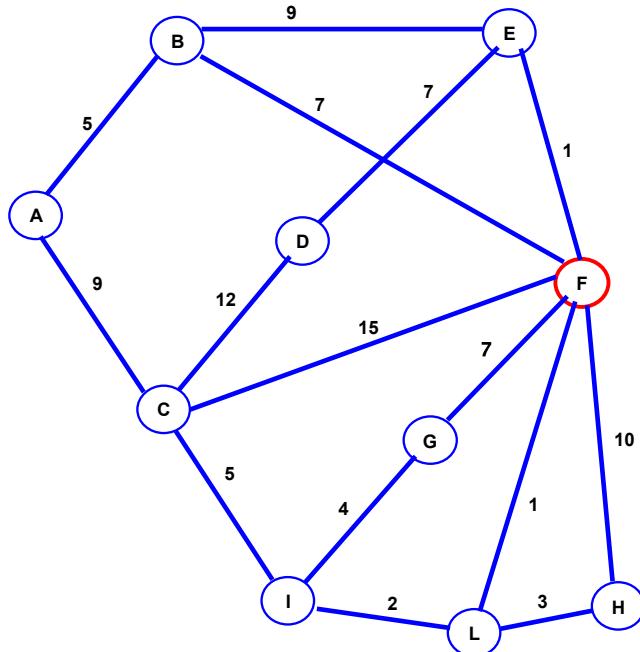
$$D = 13+7 = 20$$



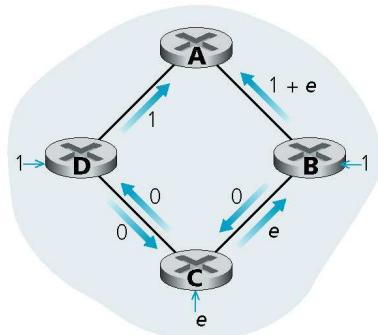
A
B = 5
C = 9
F = 12
L = 13
E = 13
I = 14
H = 16
G = 18
D = 20



A	
B = 5	AB
C = 9	AC
F = 12	ABF
L = 13	ABFL
E = 13	ABFE
I = 14	ACI
H = 16	ABFLH
G = 18	ACIG
D = 20	ABFED



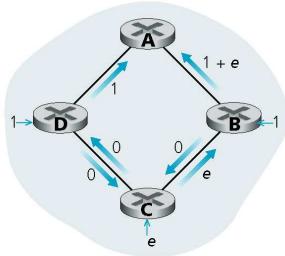
Nel LSR ogni router informa periodicamente i propri vicini sulla sua situazione dei link.



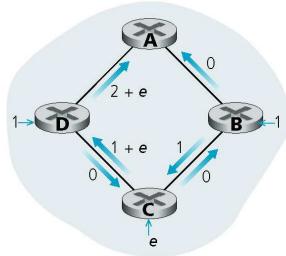
a. Routing iniziale

Nella figura i nodi B, C e D ricevono traffico dall'esterno diretto verso il nodo A.

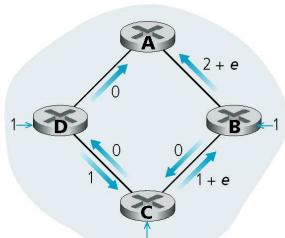
Link State Routing



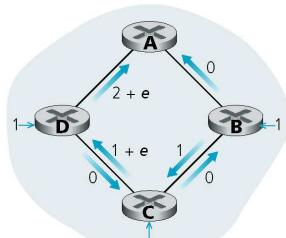
a. Routing iniziale



b. B e C rilevano il miglior percorso verso A, senso orario



c. B, C e D rilevano il miglior percorso verso A, senso antiorario



d. B, C e D, rilevano il miglior percorso verso A, senso orario

DV	LSR
Decentralizzato	Globale
Messaggi solo per i vicini	Messaggi in Broadcast
Informazioni riguardanti tutte le destinazioni	Informazioni riguardanti solo i propri link
Conteggio all'infinito	Problemi di sincronia (oscillazioni)
Messaggi solo per variazioni sui link	Maggiore traffico di messaggi (invii periodici)
Poco robusto nei confronti di nodi maliziosi	Robusto agli attacchi

Principi base:

- routing dinamico
- Intra-AS e Inter-AS routing

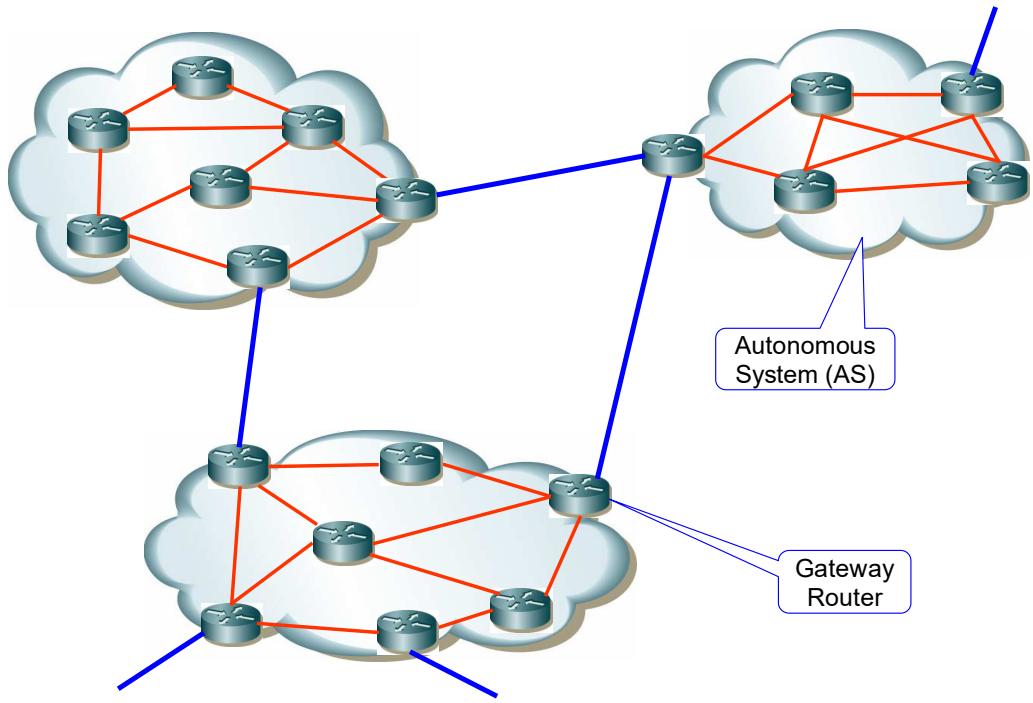
Internet è organizzata in “**autonomous systems**” (AS).

AS definition: unit of routing policy, either a single network or a group of networks that is controlled by a common network administrator on behalf of a single administrative entity (such as a university, a business enterprise, or a business division)

Ogni AS è internamente fortemente connessa.

Protocolli:

- **Interior Gateway Protocols** (IGPs) all'interno di AS (RIP, OSPF, HELLO, IS-IS)
- **Exterior Gateway Protocols** (EGPs) tra AS (EGP, BGP-4)



Routing Information Protocol:

- RIP v1 (RFC 1058)
- RIP v2 (RFC 1723, 2453).



- Basato su **Distance Vector**.
- Metrica utilizzata: **HOP** (numero di salti effettuati)

Il numero **massimo di HOP** ammesso è **15**: con reti troppo estese la convergenza risulterebbe lenta.

Oltre tale valore il router è considerato irraggiungibile.

Le **Routing Table** (RT) vengono scambiate ogni 30 secondi.

Se un percorso per un router **non viene aggiornato** entro 180 secondi, la sua distanza è posta a infinito.

Trascorsi altri 120 secondi (**garbage-collection timer**) il router viene eliminato dalla RT.

RIP v1 utilizza due tipi di messaggi :

- **REQUEST**: per chiedere informazioni ai nodi adiacenti
- **RESPONSE**: per inviare informazioni di routing

Una **tabella di routing** contiene:

- Indirizzo di destinazione
- Distanza dalla destinazione (in hop)
- Next hop: router adiacente a cui inviare i pacchetti
- Timeout
- Garbage-collection timer

Lunghezza variabile fino a 512 byte (max 25 reti di destinazione).

Command	Version	Must Be Zero
Address Family Identifier		Must Be Zero
IP Address		
Must Be Zero		
Must Be Zero		
Metric		

Command : 1 = richiesta, 2= aggiornamento.

Version : versione del protocollo.

Address Family Identifier : sempre 2 per il protocollo IP

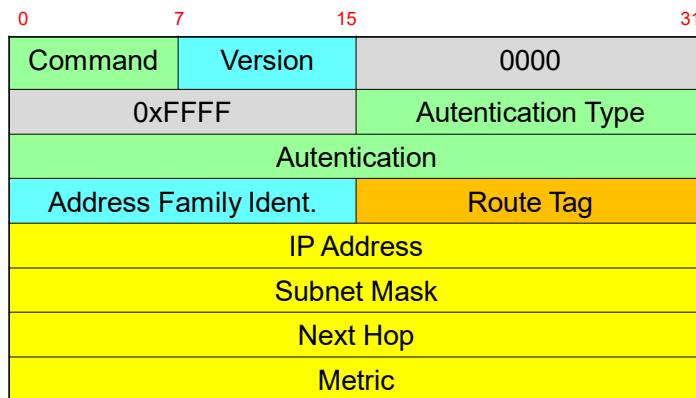
IP Address : indirizzo di destinazione (rete o sottorete)

Metric : hop count (valore compreso tra 1 e 15)

Caratteristiche di RIP v2

- Indirizzamento **CIDR** e **VLSM**
- **Autenticazione** dei messaggi
- Specifica del **next hop**
- **Split horizon, poison reverse**

Datagrammi RIP v2



Command : 1 = richiesta, 2= aggiornamento.

Version : versione del protocollo.

Address Family Identifier : sempre 2 per il protocollo IP

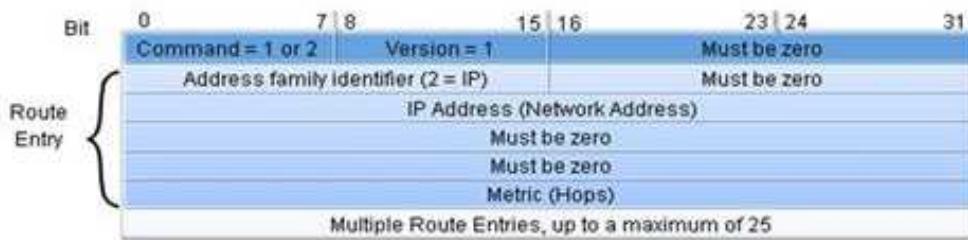
Route Tag: identificazione di route esterni

IP Address : indirizzo di destinazione (rete o sottorete)

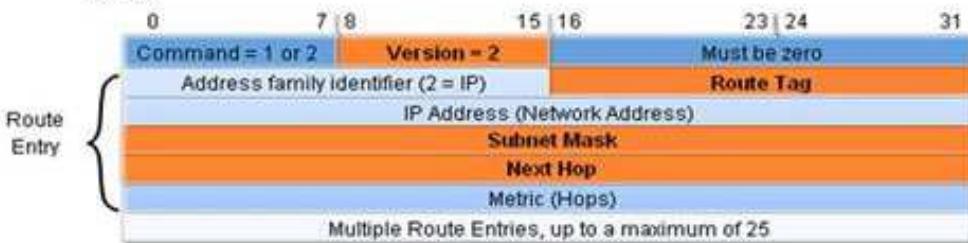
Metric : hop count (valore compreso tra 1 e 15)

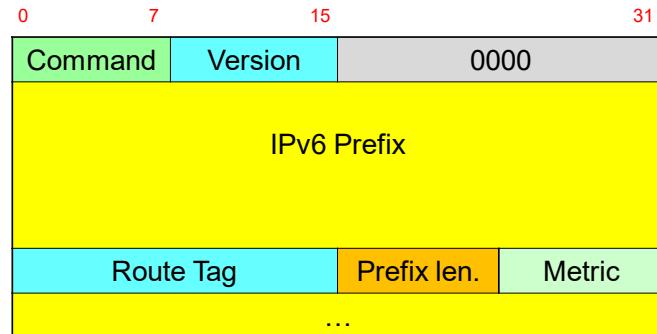
Datagrammi RIP v2

RIPv1



RIPv2





RIPng è basato su RIPv2, ma non ne rappresenta un'estensione.
È pensato esclusivamente per IPv6 (non supporta IPv4).

RIPng ha le stesse caratteristiche di RIPv2 tranne l'autenticazione

Open Shortest Path First

- OSPFv1 (RFC 1131);
- OSPFv2 (RFC 2178, 2328).

- Pensato per sostituire RIP
- Si basa sul Link State Routing
- Va bene per reti di dimensioni grandi
- I messaggi sono autenticati
- Consente l'utilizzo di metriche differenti

Il protocollo OSPF si compone di tre parti:

- **HELLO**: scoperta e verifica dei vicini
- **EXCHANGE**: sincronizzazione iniziale del DB
- **FLOODING**: aggiornamento del DB

Ogni router ha un database composto da “**link state record**”

I vari **DB** vengono aggiornati e sincronizzati tramite i “**Link State Advertisement**” (**LSA**).

Gli **LSA** sono emessi:

- quando un router riscontra un nuovo router adiacente
- quando un router – link si guasta
- quando il costo di un link cambia
- periodicamente, tipicamente ogni 30 minuti

LS Age	Options	Ls Type
Link State ID		
Advertisement Router		
Link State Sequence Number		
Link State Checksum	Lenght	
LS Data		

Version	Type	Message Length
Router ID		
Area ID		
Checksum	Autentication Type	
Autentication Data		
Rest of the OSPF Message		

Type:

1. Hello
2. DB Description Router ID: router che ha generato il messaggio
3. LS Request
4. LS Update AreaID: Ip dell'area a cui si riferisce il messaggio
5. LS ACK

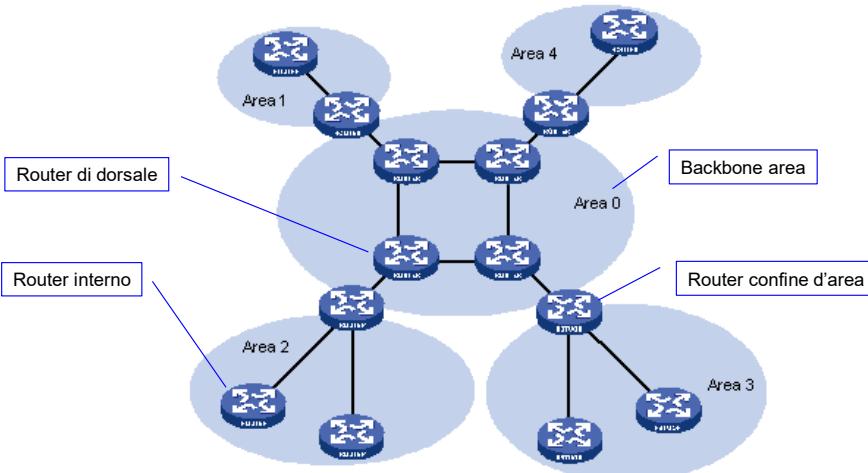
Hello: Trasmissione periodica per la scoperta dei vicini

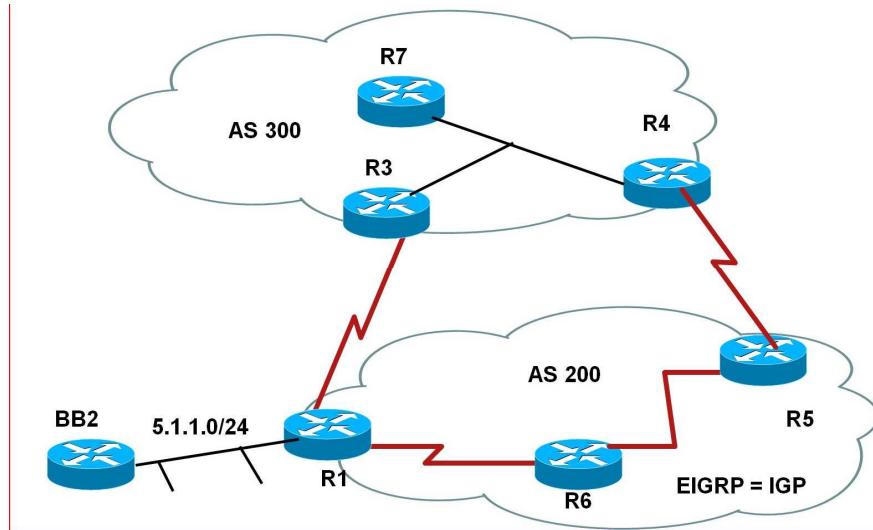
Database Description: usato per lo scambio dell'informazione dei Link-State di ogni router

LS Request: richiesta di parti specifiche del database di Link-State di un vicino

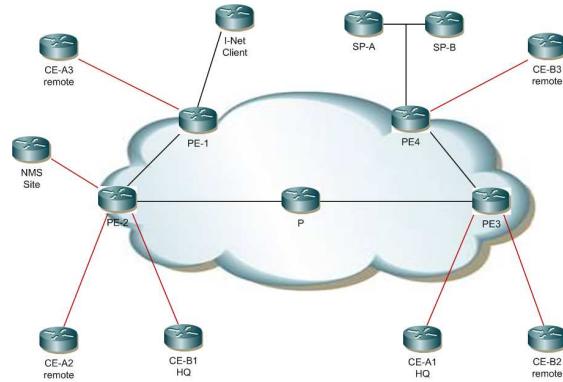
Link-State Update: trasferisce annunci di link-state ai vicini

Link-State Acknowledgments: invia una conferma di ricezione di un update di link-state





BGP



- **AS Border Router (ASBR):** router connesso ad altri sistemi autonomi
- **BGP speaker:** router che supporta il protocollo BGP (un BGP speaker non necessariamente coincide con un AS border router)
- **BGP Neighbors:** coppia di BGP speaker che si scambiano informazioni di instradamento inter-AS
 - interni: se appartengono allo stesso AS
 - esterni: se appartengono ad AS diversi